

end-to-end-bulldozer-price-regression

February 22, 2026

1 Predicting the Sale Price of Bulldozers using Machine Learning

In this notebook, we're going to go through an example machine learning project with the goal of predicting the sale price of bulldozers

1.1 1. Problem Definition

How well can we predict the future sale price of a bulldozer, give it's characteristics and previous examples of how much similar bulldozers have been sold for?

1.2 2. Data

The data is downloaded from the Kaggle Bluebook for bulldozers competition:
<https://www.kaggle.com/competitions/bluebook-for-bulldozers/data>

The data for this competition is split into three parts:

- Train.csv is the training set, which contains data through the end of 2011.
- Valid.csv is the validation set, which contains data from January 1, 2012 - April 30, 2012 You make predictions on this set throughout the majority of the competition. Your score on this set is used to create the public leaderboard.
- Test.csv is the test set, which won't be released until the last week of the competition. It contains data from May 1, 2012 - November 2012. Your score on the test set determines your final rank for the competition.

1.3 3. Evaluation

The evaluation metric for this competition is the RMSLE (root mean squared log error) between the actual and predicted auction prices.

For more on the evaluation of this project check: <https://www.kaggle.com/competitions/bluebook-for-bulldozers/overview>

Note: The goal for most regression evaluation metrics is to minimize the error. For example, our goal for this project will be to build a machine learning model which minimises RSMLE.

1.4 4.Features

Kaggle provides a data dictionary: <https://www.kaggle.com/competitions/bluebook-for-bulldozers/data>

```
[131]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import sklearn
```

```
[132]: # Import the data training & validation

df = pd.read_csv("./data/TrainAndValid.csv", low_memory=False)
```

```
[133]: df.info()
```

<class 'pandas.core.frame.DataFrame'>			
RangeIndex: 412698 entries, 0 to 412697			
Data columns (total 53 columns):			
#	Column	Non-Null Count	Dtype
---	---	-----	----
0	SalesID	412698 non-null	int64
1	SalePrice	412698 non-null	float64
2	MachineID	412698 non-null	int64
3	ModelID	412698 non-null	int64
4	datasource	412698 non-null	int64
5	auctioneerID	392562 non-null	float64
6	YearMade	412698 non-null	int64
7	MachineHoursCurrentMeter	147504 non-null	float64
8	UsageBand	73670 non-null	object
9	saledate	412698 non-null	object
10	fiModelDesc	412698 non-null	object
11	fiBaseModel	412698 non-null	object
12	fiSecondaryDesc	271971 non-null	object
13	fiModelSeries	58667 non-null	object
14	fiModelDescriptor	74816 non-null	object
15	ProductSize	196093 non-null	object
16	fiProductClassDesc	412698 non-null	object
17	state	412698 non-null	object
18	ProductGroup	412698 non-null	object
19	ProductGroupDesc	412698 non-null	object
20	Drive_System	107087 non-null	object
21	Enclosure	412364 non-null	object
22	Forks	197715 non-null	object
23	Pad_Type	81096 non-null	object
24	Ride_Control	152728 non-null	object
25	Stick	81096 non-null	object
26	Transmission	188007 non-null	object
27	Turbocharged	81096 non-null	object
28	Blade_Extension	25983 non-null	object
29	Blade_Width	25983 non-null	object
30	Enclosure_Type	25983 non-null	object
31	Engine_Horsepower	25983 non-null	object

```

32 Hydraulics           330133 non-null  object
33 Pushblock            25983 non-null   object
34 Ripper               106945 non-null  object
35 Scarifier            25994 non-null   object
36 Tip_Control          25983 non-null   object
37 Tire_Size             97638 non-null   object
38 Coupler              220679 non-null  object
39 Coupler_System        44974 non-null   object
40 Grouser_Tracks       44875 non-null   object
41 Hydraulics_Flow      44875 non-null   object
42 Track_Type           102193 non-null  object
43 Undercarriage_Pad_Width 102916 non-null  object
44 Stick_Length          102261 non-null  object
45 Thumb                 102332 non-null  object
46 Pattern_Changer       102261 non-null  object
47 Grouser_Type          102193 non-null  object
48 Backhoe_Mounting     80712 non-null   object
49 Blade_Type            81875 non-null   object
50 Travel_Controls        81877 non-null  object
51 Differential_Type     71564 non-null   object
52 Steering_Controls     71522 non-null   object
dtypes: float64(3), int64(5), object(45)
memory usage: 166.9+ MB

```

[134]: df.isna().sum()

```

[134]: SalesID                  0
SalePrice                 0
MachineID                0
ModelID                  0
datasource                0
auctioneeringID          20136
YearMade                  0
MachineHoursCurrentMeter 265194
UsageBand                 339028
saledate                  0
fiModelDesc               0
fiBaseModel               0
fiSecondaryDesc           140727
fiModelSeries              354031
fiModelDescriptor          337882
ProductSize                216605
fiProductClassDesc         0
state                      0
ProductGroup              0
ProductGroupDesc           0
Drive_System               305611

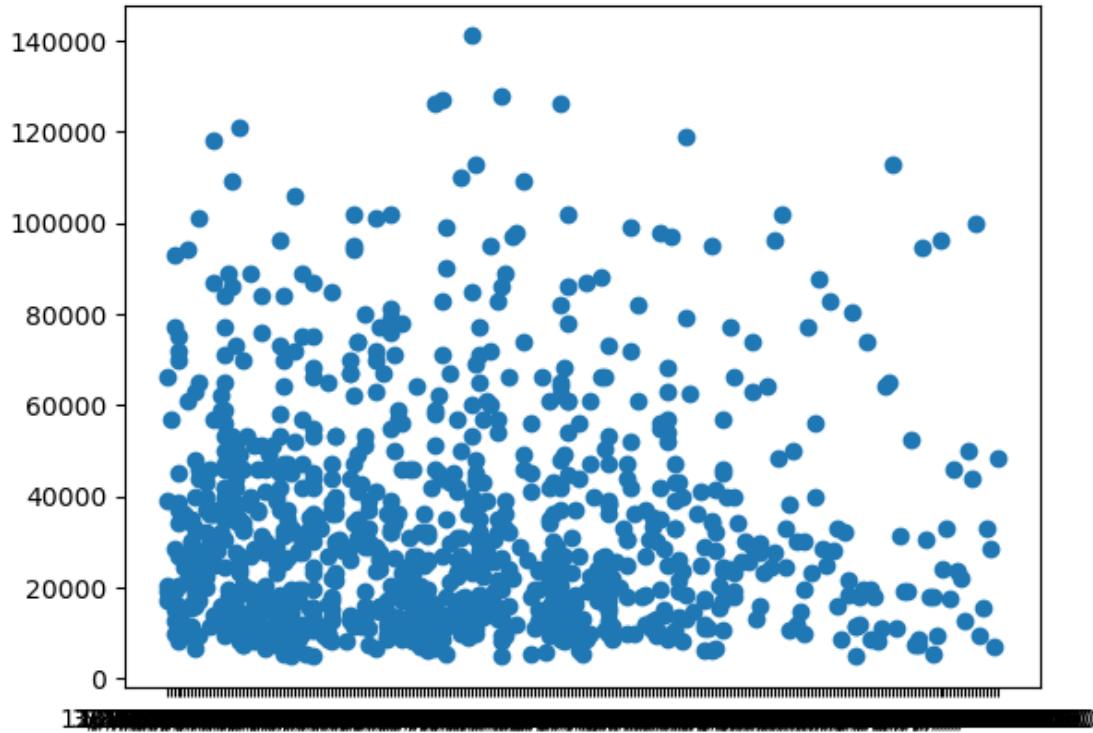
```

```
Enclosure           334
Forks              214983
Pad_Type           331602
Ride_Control       259970
Stick              331602
Transmission       224691
Turbocharged       331602
Blade_Extension    386715
Blade_Width         386715
Enclosure_Type     386715
Engine_Horsepower  386715
Hydraulics          82565
Pushblock           386715
Ripper              305753
Scarifier            386704
Tip_Control         386715
Tire_Size            315060
Coupler              192019
Coupler_System      367724
Grouser_Tracks     367823
Hydraulics_Flow     367823
Track_Type           310505
Undercarriage_Pad_Width 309782
Stick_Length         310437
Thumb                310366
Pattern_Changer      310437
Grouser_Type         310505
Backhoe_Mounting     331986
Blade_Type            330823
Travel_Controls       330821
Differential_Type    341134
Steering_Controls     341176
dtype: int64
```

```
[135]: fig, ax = plt.subplots()

ax.scatter(df["saledate"][:1000], df["SalePrice"][:1000])
```

```
[135]: <matplotlib.collections.PathCollection at 0x235bf3c9950>
```



```
[136]: df.saledate[:1000]
```

```
[136]: 0      11/16/2006 0:00
       1      3/26/2004 0:00
       2      2/26/2004 0:00
       3      5/19/2011 0:00
       4      7/23/2009 0:00
       ...
      995     7/16/2009 0:00
      996     6/14/2007 0:00
      997     9/22/2005 0:00
      998     7/28/2005 0:00
      999     6/16/2011 0:00
Name: saledate, Length: 1000, dtype: object
```

```
[137]: #df["SalePrice"].hist() # taking too long
```

1.5 Parsing Dates

When we work with time series data, we want to enrich the time & date component as much as possible.

We can do that by telling pandas which of our columns has dates in it using the `parse_dates` parameter

```
[138]: # Import data again but this time parse dates
df = pd.read_csv("./data/TrainAndValid.csv", low_memory=False, parse_dates=[-1]
                 =[["saledate"]])
```

```
[139]: df.saledate.dtype
```

```
[139]: dtype('M8[ns]')
```

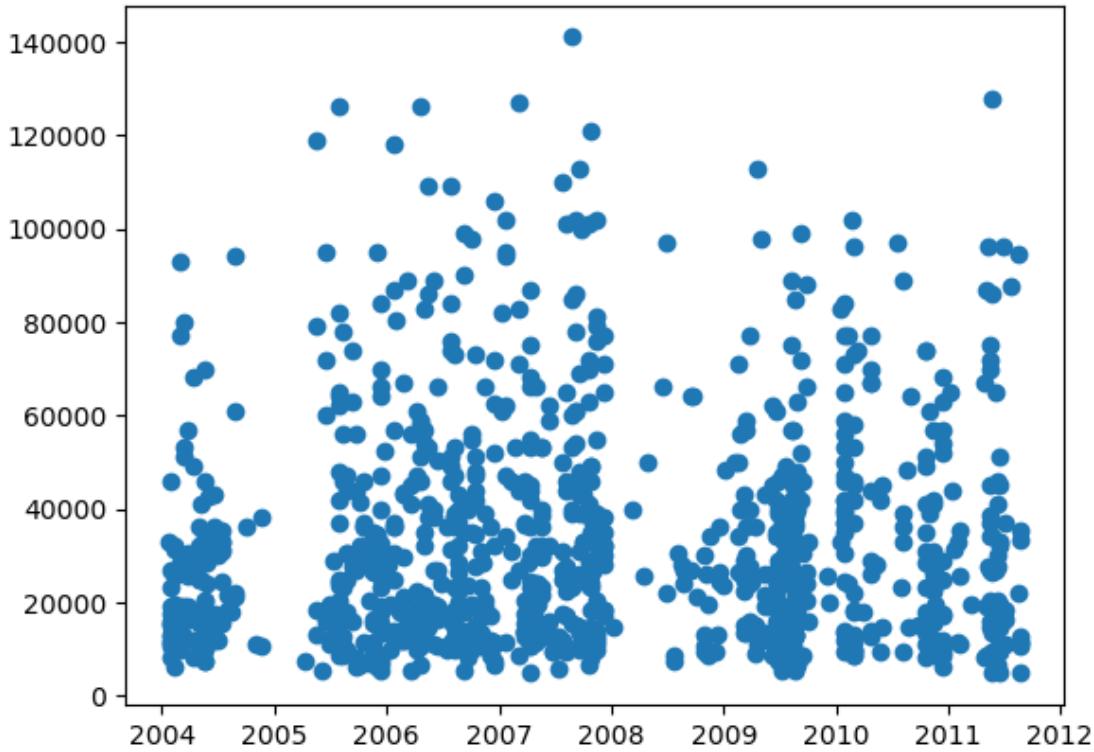
```
[140]: df.saledate[:1000]
```

```
[140]: 0      2006-11-16
       1      2004-03-26
       2      2004-02-26
       3      2011-05-19
       4      2009-07-23
       ...
      995    2009-07-16
      996    2007-06-14
      997    2005-09-22
      998    2005-07-28
      999    2011-06-16
Name: saledate, Length: 1000, dtype: datetime64[ns]
```

```
[141]: fig, ax = plt.subplots()
```

```
ax.scatter(df[["saledate"]][:1000], df[["SalePrice"]][:1000])
```

```
[141]: <matplotlib.collections.PathCollection at 0x235c53879d0>
```



[142]: df.head()

```

[142]:   SalesID  SalePrice  MachineID  ModelID  datasource  auctioneerID  YearMade \
0    1139246     66000.0    999089      3157        121         3.0    2004
1    1139248     57000.0    117657       77        121         3.0    1996
2    1139249     10000.0    434808      7009        121         3.0    2001
3    1139251     38500.0   1026470      332        121         3.0    2001
4    1139253     11000.0   1057373     17311        121         3.0    2007

                           MachineHoursCurrentMeter UsageBand  saledate ... Undercarriage_Pad_Width \
0                            68.0          Low 2006-11-16 ...                   NaN
1                            4640.0         Low 2004-03-26 ...                   NaN
2                            2838.0        High 2004-02-26 ...                   NaN
3                            3486.0        High 2011-05-19 ...                   NaN
4                            722.0       Medium 2009-07-23 ...                   NaN

   Stick_Length  Thumb  Pattern_Changer  Grouser_Type  Backhoe_Mounting  Blade_Type \
0        NaN     NaN           NaN          NaN          NaN          NaN
1        NaN     NaN           NaN          NaN          NaN          NaN
2        NaN     NaN           NaN          NaN          NaN          NaN
3        NaN     NaN           NaN          NaN          NaN          NaN
4        NaN     NaN           NaN          NaN          NaN          NaN

```

```

Travel_Controls Differential_Type Steering_Controls
0           NaN      Standard    Conventional
1           NaN      Standard    Conventional
2           NaN          NaN        NaN
3           NaN          NaN        NaN
4           NaN          NaN        NaN

```

[5 rows x 53 columns]

[143]: df.head().T

[143]:

	0 \
SalesID	1139246
SalePrice	66000.0
MachineID	999089
ModelID	3157
datasource	121
auctioneerID	3.0
YearMade	2004
MachineHoursCurrentMeter	68.0
UsageBand	Low
saledate	2006-11-16 00:00:00
fiModelDesc	521D
fiBaseModel	521
fiSecondaryDesc	D
fiModelSeries	NaN
fiModelDescriptor	NaN
ProductSize	NaN
fiProductClassDesc	Wheel Loader - 110.0 to 120.0 Horsepower
state	Alabama
ProductGroup	WL
ProductGroupDesc	Wheel Loader
Drive_System	NaN
Enclosure	EROPS w AC
Forks	None or Unspecified
Pad_Type	NaN
Ride_Control	None or Unspecified
Stick	NaN
Transmission	NaN
Turbocharged	NaN
Blade_Extension	NaN
Blade_Width	NaN
Enclosure_Type	NaN
Engine_Horsepower	NaN
Hydraulics	2 Valve
Pushblock	NaN

Ripper		NaN
Scarifier		NaN
Tip_Control		NaN
Tire_Size	None or Unspecified	
Coupler	None or Unspecified	
Coupler_System		NaN
Grouser_Tracks		NaN
Hydraulics_Flow		NaN
Track_Type		NaN
Undercarriage_Pad_Width		NaN
Stick_Length		NaN
Thumb		NaN
Pattern_Changer		NaN
Grouser_Type		NaN
Backhoe_Mounting		NaN
Blade_Type		NaN
Travel_Controls		NaN
Differential_Type	Standard	
Steering_Controls	Conventional	
		1 \
SalesID	1139248	
SalePrice	57000.0	
MachineID	117657	
ModelID	77	
datasource	121	
auctioneerID	3.0	
YearMade	1996	
MachineHoursCurrentMeter	4640.0	
UsageBand	Low	
saledate	2004-03-26 00:00:00	
fiModelDesc	950FII	
fiBaseModel	950	
fiSecondaryDesc	F	
fiModelSeries	II	
fiModelDescriptor	NaN	
ProductSize	Medium	
fiProductClassDesc	Wheel Loader - 150.0 to 175.0 Horsepower	
state	North Carolina	
ProductGroup	WL	
ProductGroupDesc	Wheel Loader	
Drive_System	NaN	
Enclosure	EROPS w AC	
Forks	None or Unspecified	
Pad_Type	NaN	
Ride_Control	None or Unspecified	
Stick	NaN	

Transmission		NaN
Turbocharged		NaN
Blade_Extension		NaN
Blade_Width		NaN
Enclosure_Type		NaN
Engine_Horsepower		NaN
Hydraulics	2	Valve
Pushblock		NaN
Ripper		NaN
Scarifier		NaN
Tip_Control		NaN
Tire_Size		23.5
Coupler		None or Unspecified
Coupler_System		NaN
Grouser_Tracks		NaN
Hydraulics_Flow		NaN
Track_Type		NaN
Undercarriage_Pad_Width		NaN
Stick_Length		NaN
Thumb		NaN
Pattern_Changer		NaN
Grouser_Type		NaN
Backhoe_Mounting		NaN
Blade_Type		NaN
Travel_Controls		NaN
Differential_Type		Standard
Steering_Controls		Conventional
	2	\
SalesID		1139249
SalePrice		10000.0
MachineID		434808
ModelID		7009
datasource		121
auctioneerID		3.0
YearMade		2001
MachineHoursCurrentMeter		2838.0
UsageBand		High
saledate		2004-02-26 00:00:00
fiModelDesc		226
fiBaseModel		226
fiSecondaryDesc		NaN
fiModelSeries		NaN
fiModelDescriptor		NaN
ProductSize		NaN
fiProductClassDesc	Skid Steer Loader - 1351.0 to 1601.0 Lb Operat...	
state		New York

ProductGroup	SSL
ProductGroupDesc	Skid Steer Loaders
Drive_System	NaN
Enclosure	OROPS
Forks	None or Unspecified
Pad_Type	NaN
Ride_Control	NaN
Stick	NaN
Transmission	NaN
Turbocharged	NaN
Blade_Extension	NaN
Blade_Width	NaN
Enclosure_Type	NaN
Engine_Horsepower	NaN
Hydraulics	Auxiliary
Pushblock	NaN
Ripper	NaN
Scarifier	NaN
Tip_Control	NaN
Tire_Size	NaN
Coupler	None or Unspecified
Coupler_System	None or Unspecified
Grouser_Tracks	None or Unspecified
Hydraulics_Flow	Standard
Track_Type	NaN
Undercarriage_Pad_Width	NaN
Stick_Length	NaN
Thumb	NaN
Pattern_Changer	NaN
Grouser_Type	NaN
Backhoe_Mounting	NaN
Blade_Type	NaN
Travel_Controls	NaN
Differential_Type	NaN
Steering_Controls	NaN
SalesID	3 \
SalePrice	1139251
MachineID	38500.0
ModelID	1026470
datasource	332
auctioneerID	121
YearMade	3.0
MachineHoursCurrentMeter	2001
UsageBand	3486.0
saledate	High
2011-05-19 00:00:00	

fiModelDesc	PC120-6E
fiBaseModel	PC120
fiSecondaryDesc	NaN
fiModelSeries	-6E
fiModelDescriptor	NaN
ProductSize	Small
fiProductClassDesc	Hydraulic Excavator, Track - 12.0 to 14.0 Metr...
state	Texas
ProductGroup	TEX
ProductGroupDesc	Track Excavators
Drive_System	NaN
Enclosure	EROPS w AC
Forks	NaN
Pad_Type	NaN
Ride_Control	NaN
Stick	NaN
Transmission	NaN
Turbocharged	NaN
Blade_Extension	NaN
Blade_Width	NaN
Enclosure_Type	NaN
Engine_Horsepower	NaN
Hydraulics	2 Valve
Pushblock	NaN
Ripper	NaN
Scarifier	NaN
Tip_Control	NaN
Tire_Size	NaN
Coupler	None or Unspecified
Coupler_System	NaN
Grouser_Tracks	NaN
Hydraulics_Flow	NaN
Track_Type	NaN
Undercarriage_Pad_Width	NaN
Stick_Length	NaN
Thumb	NaN
Pattern_Changer	NaN
Grouser_Type	NaN
Backhoe_Mounting	NaN
Blade_Type	NaN
Travel_Controls	NaN
Differential_Type	NaN
Steering_Controls	NaN
	4
SalesID	1139253
SalePrice	11000.0

MachineID		1057373
ModelID		17311
datasource		121
auctioneerID		3.0
YearMade		2007
MachineHoursCurrentMeter		722.0
UsageBand		Medium
saledate		2009-07-23 00:00:00
fiModelDesc		S175
fiBaseModel		S175
fiSecondaryDesc		NaN
fiModelSeries		NaN
fiModelDescriptor		NaN
ProductSize		NaN
fiProductClassDesc	Skid Steer Loader - 1601.0 to 1751.0 Lb Operat...	
state	New York	
ProductGroup	SSL	
ProductGroupDesc	Skid Steer Loaders	
Drive_System	NaN	
Enclosure	EROPS	
Forks	None or Unspecified	
Pad_Type	NaN	
Ride_Control	NaN	
Stick	NaN	
Transmission	NaN	
Turbocharged	NaN	
Blade_Extension	NaN	
Blade_Width	NaN	
Enclosure_Type	NaN	
Engine_Horsepower	NaN	
Hydraulics	Auxiliary	
Pushblock	NaN	
Ripper	NaN	
Scarifier	NaN	
Tip_Control	NaN	
Tire_Size	NaN	
Coupler	None or Unspecified	
Coupler_System	None or Unspecified	
Grouser_Tracks	None or Unspecified	
Hydraulics_Flow	Standard	
Track_Type	NaN	
Undercarriage_Pad_Width	NaN	
Stick_Length	NaN	
Thumb	NaN	
Pattern_Changer	NaN	
Grouser_Type	NaN	
Backhoe_Mounting	NaN	

```
Blade_Type           NaN
Travel_Controls      NaN
Differential_Type   NaN
Steering_Controls   NaN
```

```
[144]: df.saledate.head(20)
```

```
[144]: 0    2006-11-16
1    2004-03-26
2    2004-02-26
3    2011-05-19
4    2009-07-23
5    2008-12-18
6    2004-08-26
7    2005-11-17
8    2009-08-27
9    2007-08-09
10   2008-08-21
11   2006-08-24
12   2005-10-20
13   2006-01-26
14   2006-01-03
15   2006-11-16
16   2007-06-14
17   2010-01-28
18   2006-03-09
19   2005-11-17
Name: saledate, dtype: datetime64[ns]
```

1.5.1 Sort DataFrame by saledate

When working with Time Series Data it is a good idea to sort it by date.

```
[145]: df.sort_values(by=["saledate"], inplace = True, ascending=True)
```

```
[146]: df.saledate.head(20)
```

```
[146]: 205615  1989-01-17
274835  1989-01-31
141296  1989-01-31
212552  1989-01-31
62755   1989-01-31
54653   1989-01-31
81383   1989-01-31
204924  1989-01-31
135376  1989-01-31
113390  1989-01-31
113394  1989-01-31
```

```

116419 1989-01-31
32138 1989-01-31
127610 1989-01-31
76171 1989-01-31
127000 1989-01-31
128130 1989-01-31
127626 1989-01-31
55455 1989-01-31
55454 1989-01-31
Name: saledate, dtype: datetime64[ns]

```

[147]: df.head()

```

[147]:      SalesID  SalePrice  MachineID  ModelID  datasource  auctioneerID \
205615    1646770     9500.0    1126363     8434        132       18.0
274835    1821514    14000.0    1194089     10150       132       99.0
141296    1505138    50000.0    1473654     4139       132       99.0
212552    1671174    16000.0    1327630     8591       132       99.0
62755     1329056    22000.0    1336053     4089       132       99.0

      YearMade  MachineHoursCurrentMeter  UsageBand  saledate  ... \
205615      1974                  NaN          NaN  1989-01-17 ...
274835      1980                  NaN          NaN  1989-01-31 ...
141296      1978                  NaN          NaN  1989-01-31 ...
212552      1980                  NaN          NaN  1989-01-31 ...
62755       1984                  NaN          NaN  1989-01-31 ...

      Undercarriage_Pad_Width  Stick_Length  Thumb  Pattern_Changer \
205615            NaN             NaN     NaN           NaN
274835            NaN             NaN     NaN           NaN
141296            NaN             NaN     NaN           NaN
212552            NaN             NaN     NaN           NaN
62755            NaN             NaN     NaN           NaN

      Grouser_Type  Backhoe_Mounting  Blade_Type  Travel_Controls \
205615            NaN   None or Unspecified  Straight  None or Unspecified
274835            NaN                 NaN        NaN           NaN
141296            NaN   None or Unspecified  Straight  None or Unspecified
212552            NaN                 NaN        NaN           NaN
62755            NaN   None or Unspecified        PAT           Lever

      Differential_Type  Steering_Controls
205615            NaN                 NaN
274835          Standard      Conventional
141296            NaN                 NaN
212552          Standard      Conventional
62755            NaN                 NaN

```

```
[5 rows x 53 columns]
```

1.5.2 Make Copy of the original DataFrame

We make a copy of the original dataframe so when we manipulate the copy, we have still got our original data

```
[148]: # Make a copy  
df_temp = df.copy()
```

```
[149]: df_temp.saledate.head(20)
```

```
[149]: 205615    1989-01-17  
274835    1989-01-31  
141296    1989-01-31  
212552    1989-01-31  
62755     1989-01-31  
54653     1989-01-31  
81383     1989-01-31  
204924    1989-01-31  
135376    1989-01-31  
113390    1989-01-31  
113394    1989-01-31  
116419    1989-01-31  
32138     1989-01-31  
127610    1989-01-31  
76171     1989-01-31  
127000    1989-01-31  
128130    1989-01-31  
127626    1989-01-31  
55455     1989-01-31  
55454     1989-01-31  
Name: saledate, dtype: datetime64[ns]
```

1.5.3 Add datetime parameters for saledate column

```
[150]: df_temp["saleYear"] = df_temp.saledate.dt.year  
df_temp["saleMonth"] = df_temp.saledate.dt.month  
df_temp["saleDay"] = df_temp.saledate.dt.day  
df_temp["saleDayOfWeek"] = df_temp.saledate.dt.dayofweek  
df_temp["saleDayOfYear"] = df_temp.saledate.dt.dayofyear
```

```
[151]: df_temp.head().T
```

```
[151]:  
SalesID          205615  \\\nSalePrice        1646770  
                 9500.0
```

MachineID		1126363
ModelID		8434
datasource		132
auctioneerID		18.0
YearMade		1974
MachineHoursCurrentMeter		Nan
UsageBand		Nan
saledate		1989-01-17 00:00:00
fiModelDesc		TD20
fiBaseModel		TD20
fiSecondaryDesc		Nan
fiModelSeries		Nan
fiModelDescriptor		Nan
ProductSize		Medium
fiProductClassDesc	Track Type Tractor, Dozer - 105.0 to 130.0 Hor...	
state		Texas
ProductGroup		TTT
ProductGroupDesc		Track Type Tractors
Drive_System		Nan
Enclosure		OROPS
Forks		Nan
Pad_Type		Nan
Ride_Control		Nan
Stick		Nan
Transmission		Direct Drive
Turbocharged		Nan
Blade_Extension		Nan
Blade_Width		Nan
Enclosure_Type		Nan
Engine_Horsepower		Nan
Hydraulics	2	Valve
Pushblock		Nan
Ripper		None or Unspecified
Scarifier		Nan
Tip_Control		Nan
Tire_Size		Nan
Coupler		Nan
Coupler_System		Nan
Grouser_Tracks		Nan
Hydraulics_Flow		Nan
Track_Type		Nan
Undercarriage_Pad_Width		Nan
Stick_Length		Nan
Thumb		Nan
Pattern_Changer		Nan
Grouser_Type		Nan
Backhoe_Mounting		None or Unspecified

Blade_Type	Straight
Travel_Controls	None or Unspecified
Differential_Type	NaN
Steering_Controls	NaN
saleYear	1989
saleMonth	1
saleDay	17
saleDayOfWeek	1
saleDayOfYear	17
	274835 \
SalesID	1821514
SalePrice	14000.0
MachineID	1194089
ModelID	10150
datasource	132
auctioneerID	99.0
YearMade	1980
MachineHoursCurrentMeter	NaN
UsageBand	NaN
saledate	1989-01-31 00:00:00
fiModelDesc	A66
fiBaseModel	A66
fiSecondaryDesc	NaN
fiModelSeries	NaN
fiModelDescriptor	NaN
ProductSize	NaN
fiProductClassDesc	Wheel Loader - 120.0 to 135.0 Horsepower
state	Florida
ProductGroup	WL
ProductGroupDesc	Wheel Loader
Drive_System	NaN
Enclosure	OROPS
Forks	None or Unspecified
Pad_Type	NaN
Ride_Control	None or Unspecified
Stick	NaN
Transmission	NaN
Turbocharged	NaN
Blade_Extension	NaN
Blade_Width	NaN
Enclosure_Type	NaN
Engine_Horsepower	NaN
Hydraulics	2 Valve
Pushblock	NaN
Ripper	NaN
Scarifier	NaN

Tip_Control		NaN
Tire_Size		None or Unspecified
Coupler		None or Unspecified
Coupler_System		NaN
Grouser_Tracks		NaN
Hydraulics_Flow		NaN
Track_Type		NaN
Undercarriage_Pad_Width		NaN
Stick_Length		NaN
Thumb		NaN
Pattern_Changer		NaN
Grouser_Type		NaN
Backhoe_Mounting		NaN
Blade_Type		NaN
Travel_Controls		NaN
Differential_Type		Standard
Steering_Controls		Conventional
saleYear		1989
saleMonth		1
saleDay		31
saleDayOfWeek		1
saleDayOfYear		31
		141296 \
SalesID		1505138
SalePrice		50000.0
MachineID		1473654
ModelID		4139
datasource		132
auctioneerID		99.0
YearMade		1978
MachineHoursCurrentMeter		NaN
UsageBand		NaN
saledate		1989-01-31 00:00:00
fiModelDesc		D7G
fiBaseModel		D7
fiSecondaryDesc		G
fiModelSeries		NaN
fiModelDescriptor		NaN
ProductSize		Large
fiProductClassDesc	Track Type Tractor, Dozer - 190.0 to 260.0 Hor...	
state		Florida
ProductGroup		TTT
ProductGroupDesc		Track Type Tractors
Drive_System		NaN
Enclosure		OROPS
Forks		NaN

Pad_Type	NaN
Ride_Control	NaN
Stick	NaN
Transmission	Standard
Turbocharged	NaN
Blade_Extension	NaN
Blade_Width	NaN
Enclosure_Type	NaN
Engine_Horsepower	NaN
Hydraulics	2 Valve
Pushblock	NaN
Ripper	None or Unspecified
Scarifier	NaN
Tip_Control	NaN
Tire_Size	NaN
Coupler	NaN
Coupler_System	NaN
Grouser_Tracks	NaN
Hydraulics_Flow	NaN
Track_Type	NaN
Undercarriage_Pad_Width	NaN
Stick_Length	NaN
Thumb	NaN
Pattern_Changer	NaN
Grouser_Type	NaN
Backhoe_Mounting	None or Unspecified
Blade_Type	Straight
Travel_Controls	None or Unspecified
Differential_Type	NaN
Steering_Controls	NaN
saleYear	1989
saleMonth	1
saleDay	31
saleDayOfWeek	1
saleDayOfYear	31
	212552 \
SalesID	1671174
SalePrice	16000.0
MachineID	1327630
ModelID	8591
datasource	132
auctioneerID	99.0
YearMade	1980
MachineHoursCurrentMeter	NaN
UsageBand	NaN
saledate	1989-01-31 00:00:00

fiModelDesc	A62
fiBaseModel	A62
fiSecondaryDesc	NaN
fiModelSeries	NaN
fiModelDescriptor	NaN
ProductSize	NaN
fiProductClassDesc	Wheel Loader - Unidentified
state	Florida
ProductGroup	WL
ProductGroupDesc	Wheel Loader
Drive_System	NaN
Enclosure	EROPS
Forks	None or Unspecified
Pad_Type	NaN
Ride_Control	None or Unspecified
Stick	NaN
Transmission	NaN
Turbocharged	NaN
Blade_Extension	NaN
Blade_Width	NaN
Enclosure_Type	NaN
Engine_Horsepower	NaN
Hydraulics	2 Valve
Pushblock	NaN
Ripper	NaN
Scarifier	NaN
Tip_Control	NaN
Tire_Size	None or Unspecified
Coupler	None or Unspecified
Coupler_System	NaN
Grouser_Tracks	NaN
Hydraulics_Flow	NaN
Track_Type	NaN
Undercarriage_Pad_Width	NaN
Stick_Length	NaN
Thumb	NaN
Pattern_Changer	NaN
Grouser_Type	NaN
Backhoe_Mounting	NaN
Blade_Type	NaN
Travel_Controls	NaN
Differential_Type	Standard
Steering_Controls	Conventional
saleYear	1989
saleMonth	1
saleDay	31
saleDayOfWeek	1

		62755
SalesID		1329056
SalePrice		22000.0
MachineID		1336053
ModelID		4089
datasource		132
auctioneerID		99.0
YearMade		1984
MachineHoursCurrentMeter		NaN
UsageBand		NaN
saledate		1989-01-31 00:00:00
fiModelDesc		D3B
fiBaseModel		D3
fiSecondaryDesc		B
fiModelSeries		NaN
fiModelDescriptor		NaN
ProductSize		NaN
fiProductClassDesc	Track Type Tractor, Dozer - 20.0 to 75.0 Horse...	
state		Florida
ProductGroup		TTT
ProductGroupDesc		Track Type Tractors
Drive_System		NaN
Enclosure		OROPS
Forks		NaN
Pad_Type		NaN
Ride_Control		NaN
Stick		NaN
Transmission		Standard
Turbocharged		NaN
Blade_Extension		NaN
Blade_Width		NaN
Enclosure_Type		NaN
Engine_Horsepower		NaN
Hydraulics		2 Valve
Pushblock		NaN
Ripper		None or Unspecified
Scarifier		NaN
Tip_Control		NaN
Tire_Size		NaN
Coupler		NaN
Coupler_System		NaN
Grouser_Tracks		NaN
Hydraulics_Flow		NaN
Track_Type		NaN
Undercarriage_Pad_Width		NaN

Stick_Length		NaN
Thumb		NaN
Pattern_Changer		NaN
Grouser_Type		NaN
Backhoe_Mounting		None or Unspecified
Blade_Type		PAT
Travel_Controls		Lever
Differential_Type		NaN
Steering_Controls		NaN
saleYear		1989
saleMonth		1
saleDay		31
saleDayOfWeek		1
saleDayOfYear		31

```
[152]: # Now we have enriched our DataFrame with date time features we can remove →saledate
        df_temp.drop("saledate",axis=1,inplace=True)
```

```
[153]: # Check the values of diffrent columns
        df_temp.state.value_counts()
```

```
[153]: state
        Florida          67320
        Texas            53110
        California       29761
        Washington      16222
        Georgia          14633
        Maryland          13322
        Mississippi      13240
        Ohio              12369
        Illinois          11540
        Colorado          11529
        New Jersey        11156
        North Carolina    10636
        Tennessee         10298
        Alabama           10292
        Pennsylvania      10234
        South Carolina    9951
        Arizona           9364
        New York          8639
        Connecticut       8276
        Minnesota         7885
        Missouri          7178
        Nevada            6932
        Louisiana          6627
        Kentucky          5351
```

Maine	5096
Indiana	4124
Arkansas	3933
New Mexico	3631
Utah	3046
Unspecified	2801
Wisconsin	2745
New Hampshire	2738
Virginia	2353
Idaho	2025
Oregon	1911
Michigan	1831
Wyoming	1672
Montana	1336
Iowa	1336
Oklahoma	1326
Nebraska	866
West Virginia	840
Kansas	667
Delaware	510
North Dakota	480
Alaska	430
Massachusetts	347
Vermont	300
South Dakota	244
Hawaii	118
Rhode Island	83
Puerto Rico	42
Washington DC	2
Name: count, dtype: int64	

1.6 5. Modelling

We've done enough EDA (we could always do more) but let's start to do some model driven EDA.

```
[154]: # Let's build a machine learning model
from sklearn.ensemble import RandomForestRegressor

model = RandomForestRegressor(n_jobs=-1,
                             random_state = 42 #same like random_seed
                            )

#model.fit(df_temp.drop("SalePrice",axis = 1), df_temp.SalePrice) # is not working because we have some features which datatypes are object
```

1.6.1 Convert string to categories

One way we can turn all of our data into numbers is by converting them into panda categories

We can check the opportunities here: <https://pandas.pydata.org/docs/reference/api/pandas.Categorical.dtype.html>

```
[155]: pd.api.types.is_string_dtype(df_temp["UsageBand"])
```

```
[155]: False
```

```
[156]: # Find the columns which contain strings
for label, content in df_temp.items():
    if pd.api.types.is_string_dtype(content):
        print(label)
```

```
fiModelDesc
fiBaseModel
fiProductClassDesc
state
ProductGroup
ProductGroupDesc
```

```
[157]: # if you are wondering what df.items does here is a example
```

```
random_dict = {"key1": "Hello",
               "key2": "World", }
```

```
[158]: for label, content in random_dict.items():
        print(label)
```

```
key1
key2
```

```
[159]: # This will turn all of the string values into category values
```

```
for label, content in df_temp.items():
    if pd.api.types.is_string_dtype(content):
        df_temp[label] = content.astype("category").cat.as_ordered()
```

```
[160]: # Wir wandeln diese Spalte(n) in den pandas-Datentyp "category" um.
# Dabei speichert pandas nicht jeden String-Wert pro Zeile, sondern:
#   1) eine feste Liste aller möglichen Kategorien (Labels) und
#   2) pro Zeile nur einen integer Code, der auf die Kategorie zeigt.
# Vorteil: weniger Speicherverbrauch und oft schnellere Operationen bei wenigen, ↴
#         häufig wiederholten Werten.
# Optional (falls gesetzt): Mit ordered=True bekommt die Kategorie eine feste ↴
#         Reihenfolge, die Sortierung/Vergleiche beeinflusst.
```

```
[161]: df_temp.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 412698 entries, 205615 to 409203
Data columns (total 57 columns):
```

#	Column	Non-Null Count	Dtype
0	SalesID	412698	non-null int64
1	SalePrice	412698	non-null float64
2	MachineID	412698	non-null int64
3	ModelID	412698	non-null int64
4	datasource	412698	non-null int64
5	auctioneerID	392562	non-null float64
6	YearMade	412698	non-null int64
7	MachineHoursCurrentMeter	147504	non-null float64
8	UsageBand	73670	non-null object
9	fiModelDesc	412698	non-null category
10	fiBaseModel	412698	non-null category
11	fiSecondaryDesc	271971	non-null object
12	fiModelSeries	58667	non-null object
13	fiModelDescriptor	74816	non-null object
14	ProductSize	196093	non-null object
15	fiProductClassDesc	412698	non-null category
16	state	412698	non-null category
17	ProductGroup	412698	non-null category
18	ProductGroupDesc	412698	non-null category
19	Drive_System	107087	non-null object
20	Enclosure	412364	non-null object
21	Forks	197715	non-null object
22	Pad_Type	81096	non-null object
23	Ride_Control	152728	non-null object
24	Stick	81096	non-null object
25	Transmission	188007	non-null object
26	Turbocharged	81096	non-null object
27	Blade_Extension	25983	non-null object
28	Blade_Width	25983	non-null object
29	Enclosure_Type	25983	non-null object
30	Engine_Horsepower	25983	non-null object
31	Hydraulics	330133	non-null object
32	Pushblock	25983	non-null object
33	Ripper	106945	non-null object
34	Scarifier	25994	non-null object
35	Tip_Control	25983	non-null object
36	Tire_Size	97638	non-null object
37	Coupler	220679	non-null object
38	Coupler_System	44974	non-null object
39	Grouser_Tracks	44875	non-null object
40	Hydraulics_Flow	44875	non-null object
41	Track_Type	102193	non-null object
42	Undercarriage_Pad_Width	102916	non-null object
43	Stick_Length	102261	non-null object
44	Thumb	102332	non-null object
45	Pattern_Changer	102261	non-null object

```
46 Grouser_Type           102193 non-null  object
47 Backhoe_Mounting       80712 non-null   object
48 Blade_Type              81875 non-null   object
49 Travel_Controls          81877 non-null   object
50 Differential_Type        71564 non-null   object
51 Steering_Controls        71522 non-null   object
52 saleYear                  412698 non-null  int32
53 saleMonth                 412698 non-null  int32
54 saleDay                   412698 non-null  int32
55 saleDayOfWeek             412698 non-null  int32
56 saleDayOfYear              412698 non-null  int32
dtypes: category(6), float64(3), int32(5), int64(5), object(38)
memory usage: 159.3+ MB
```

```
[162]: df_temp.state.cat.categories
```

```
[162]: Index(['Alabama', 'Alaska', 'Arizona', 'Arkansas', 'California', 'Colorado',
   'Connecticut', 'Delaware', 'Florida', 'Georgia', 'Hawaii', 'Idaho',
   'Illinois', 'Indiana', 'Iowa', 'Kansas', 'Kentucky', 'Louisiana',
   'Maine', 'Maryland', 'Massachusetts', 'Michigan', 'Minnesota',
   'Mississippi', 'Missouri', 'Montana', 'Nebraska', 'Nevada',
   'New Hampshire', 'New Jersey', 'New Mexico', 'New York',
   'North Carolina', 'North Dakota', 'Ohio', 'Oklahoma', 'Oregon',
   'Pennsylvania', 'Puerto Rico', 'Rhode Island', 'South Carolina',
   'South Dakota', 'Tennessee', 'Texas', 'Unspecified', 'Utah', 'Vermont',
   'Virginia', 'Washington', 'Washington DC', 'West Virginia', 'Wisconsin',
   'Wyoming'],
  dtype='object')
```

```
[163]: df_temp.state.value_counts()
```

```
[163]: state
Florida            67320
Texas              53110
California         29761
Washington          16222
Georgia             14633
Maryland             13322
Mississippi         13240
Ohio                12369
Illinois             11540
Colorado             11529
New Jersey           11156
North Carolina       10636
Tennessee             10298
Alabama              10292
Pennsylvania          10234
South Carolina        9951
```

```

Arizona          9364
New York        8639
Connecticut     8276
Minnesota       7885
Missouri        7178
Nevada          6932
Louisiana       6627
Kentucky         5351
Maine           5096
Indiana          4124
Arkansas         3933
New Mexico       3631
Utah            3046
Unspecified      2801
Wisconsin        2745
New Hampshire    2738
Virginia         2353
Idaho           2025
Oregon          1911
Michigan         1831
Wyoming          1672
Montana          1336
Iowa            1336
Oklahoma         1326
Nebraska         866
West Virginia    840
Kansas           667
Delaware          510
North Dakota     480
Alaska            430
Massachusetts    347
Vermont          300
South Dakota     244
Hawaii            118
Rhode Island      83
Puerto Rico       42
Washington DC      2
Name: count, dtype: int64

```

[164]: # This will turn all of the object values into category values

```

for label, content in df_temp.items():
    if pd.api.types.is_object_dtype(content):
        df_temp[label] = content.astype("category").cat.as_ordered()

```

[165]: # Wir wandeln diese Spalte(n) in den pandas-Datentyp "category" um.
Dabei speichert pandas nicht jeden String-Wert pro Zeile, sondern:

```

# 1) eine feste Liste aller möglichen Kategorien (Labels) und
# 2) pro Zeile nur einen integer Code, der auf die Kategorie zeigt.
# Vorteil: weniger Speicherverbrauch und oft schnellere Operationen bei wenigen, ↴
    ↴ häufig wiederholten Werten.
# Optional (falls gesetzt): Mit ordered=True bekommt die Kategorie eine feste ↴
    ↴ Reihenfolge, die Sortierung/Vergleiche beeinflusst.

```

[166]: df_temp.info()

```

<class 'pandas.core.frame.DataFrame'>
Index: 412698 entries, 205615 to 409203
Data columns (total 57 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   SalesID          412698 non-null   int64  
 1   SalePrice         412698 non-null   float64 
 2   MachineID         412698 non-null   int64  
 3   ModelID           412698 non-null   int64  
 4   datasource        412698 non-null   int64  
 5   auctioneerID      392562 non-null   float64 
 6   YearMade          412698 non-null   int64  
 7   MachineHoursCurrentMeter 147504 non-null   float64 
 8   UsageBand         73670 non-null    category 
 9   fiModelDesc       412698 non-null   category 
 10  fiBaseModel       412698 non-null   category 
 11  fiSecondaryDesc  271971 non-null   category 
 12  fiModelSeries     58667 non-null    category 
 13  fiModelDescriptor 74816 non-null   category 
 14  ProductSize       196093 non-null   category 
 15  fiProductClassDesc 412698 non-null   category 
 16  state             412698 non-null   category 
 17  ProductGroup      412698 non-null   category 
 18  ProductGroupDesc  412698 non-null   category 
 19  Drive_System      107087 non-null   category 
 20  Enclosure         412364 non-null   category 
 21  Forks              197715 non-null   category 
 22  Pad_Type           81096 non-null    category 
 23  Ride_Control       152728 non-null   category 
 24  Stick               81096 non-null   category 
 25  Transmission        188007 non-null   category 
 26  Turbocharged        81096 non-null   category 
 27  Blade_Extension    25983 non-null    category 
 28  Blade_Width         25983 non-null    category 
 29  Enclosure_Type     25983 non-null    category 
 30  Engine_Horsepower  25983 non-null    category 
 31  Hydraulics          330133 non-null   category 
 32  Pushblock           25983 non-null    category 
 33  Ripper              106945 non-null   category 

```

```
34 Scarifier           25994 non-null  category
35 Tip_Control         25983 non-null  category
36 Tire_Size            97638 non-null  category
37 Coupler              220679 non-null category
38 Coupler_System       44974 non-null  category
39 Grouser_Tracks      44875 non-null  category
40 Hydraulics_Flow     44875 non-null  category
41 Track_Type           102193 non-null category
42 Undercarriage_Pad_Width 102916 non-null category
43 Stick_Length          102261 non-null category
44 Thumb                 102332 non-null category
45 Pattern_Changer       102261 non-null category
46 Grouser_Type          102193 non-null category
47 Backhoe_Mounting      80712 non-null  category
48 Blade_Type             81875 non-null  category
49 Travel_Controls        81877 non-null  category
50 Differential_Type      71564 non-null  category
51 Steering_Controls      71522 non-null  category
52 saleYear               412698 non-null int32
53 saleMonth              412698 non-null int32
54 saleDay                412698 non-null int32
55 saleDayOfWeek          412698 non-null int32
56 saleDayOfYear          412698 non-null int32
dtypes: category(44), float64(3), int32(5), int64(5)
memory usage: 55.4 MB
```

```
[167]: df_temp.state.cat.codes
```

```
205615    43
274835     8
141296     8
212552     8
62755      8
...
410879     4
412476     4
411927     4
407124     4
409203     4
Length: 412698, dtype: int8
```

Thanks to pandas Categories we now have a way to access all of our data in the form of numbers.

But we still have a bunch of missing data

```
[168]: df_temp.isnull().sum()/len(df_temp)
```

[168]:	SalesID	0.000000
	SalePrice	0.000000
	MachineID	0.000000
	ModelID	0.000000
	datasource	0.000000
	auctioneerID	0.048791
	YearMade	0.000000
	MachineHoursCurrentMeter	0.642586
	UsageBand	0.821492
	fiModelDesc	0.000000
	fiBaseModel	0.000000
	fiSecondaryDesc	0.340993
	fiModelSeries	0.857845
	fiModelDescriptor	0.818715
	ProductSize	0.524851
	fiProductClassDesc	0.000000
	state	0.000000
	ProductGroup	0.000000
	ProductGroupDesc	0.000000
	Drive_System	0.740520
	Enclosure	0.000809
	Forks	0.520921
	Pad_Type	0.803498
	Ride_Control	0.629928
	Stick	0.803498
	Transmission	0.544444
	Turbocharged	0.803498
	Blade_Extension	0.937041
	Blade_Width	0.937041
	Enclosure_Type	0.937041
	Engine_Horsepower	0.937041
	Hydraulics	0.200062
	Pushblock	0.937041
	Ripper	0.740864
	Scarifier	0.937014
	Tip_Control	0.937041
	Tire_Size	0.763415
	Coupler	0.465277
	Coupler_System	0.891024
	Grouser_Tracks	0.891264
	Hydraulics_Flow	0.891264
	Track_Type	0.752378
	Undercarriage_Pad_Width	0.750626
	Stick_Length	0.752213
	Thumb	0.752041
	Pattern_Changer	0.752213
	Grouser_Type	0.752378

```

Backhoe_Mounting          0.804428
Blade_Type                0.801610
Travel_Controls            0.801606
Differential_Type          0.826595
Steering_Controls          0.826697
saleYear                  0.000000
saleMonth                 0.000000
saleDay                   0.000000
saleDayOfWeek              0.000000
saleDayOfYear               0.000000
dtype: float64

```

1.6.2 Save preprocessed data

```
[169]: #Export current temp df
df_temp.to_csv("data/train_temp.csv", index = False)
```



```
[170]: # Import preprocessed data
df_temp=pd.read_csv("data/train_temp.csv", low_memory= False)
df_temp.head().T
```



```
[170]:
```

	0 \
SalesID	1646770
SalePrice	9500.0
MachineID	1126363
ModelID	8434
datasource	132
auctioneerID	18.0
YearMade	1974
MachineHoursCurrentMeter	NaN
UsageBand	NaN
fiModelDesc	TD20
fiBaseModel	TD20
fiSecondaryDesc	NaN
fiModelSeries	NaN
fiModelDescriptor	NaN
ProductSize	Medium
fiProductClassDesc	Track Type Tractor, Dozer - 105.0 to 130.0 Hor...
state	Texas
ProductGroup	TTT
ProductGroupDesc	Track Type Tractors
Drive_System	NaN
Enclosure	OROPS
Forks	NaN

Pad_Type	NaN
Ride_Control	NaN
Stick	NaN
Transmission	Direct Drive
Turbocharged	NaN
Blade_Extension	NaN
Blade_Width	NaN
Enclosure_Type	NaN
Engine_Horsepower	NaN
Hydraulics	2 Valve
Pushblock	NaN
Ripper	None or Unspecified
Scarifier	NaN
Tip_Control	NaN
Tire_Size	NaN
Coupler	NaN
Coupler_System	NaN
Grouser_Tracks	NaN
Hydraulics_Flow	NaN
Track_Type	NaN
Undercarriage_Pad_Width	NaN
Stick_Length	NaN
Thumb	NaN
Pattern_Changer	NaN
Grouser_Type	NaN
Backhoe_Mounting	None or Unspecified
Blade_Type	Straight
Travel_Controls	None or Unspecified
Differential_Type	NaN
Steering_Controls	NaN
saleYear	1989
saleMonth	1
saleDay	17
saleDayOfWeek	1
saleDayOfYear	17
	1 \
SalesID	1821514
SalePrice	14000.0
MachineID	1194089
ModelID	10150
datasource	132
auctioneerID	99.0
YearMade	1980
MachineHoursCurrentMeter	NaN
UsageBand	NaN
fiModelDesc	A66

fiBaseModel	A66
fiSecondaryDesc	NaN
fiModelSeries	NaN
fiModelDescriptor	NaN
ProductSize	NaN
fiProductClassDesc	Wheel Loader - 120.0 to 135.0 Horsepower
state	Florida
ProductGroup	WL
ProductGroupDesc	Wheel Loader
Drive_System	NaN
Enclosure	OROPS
Forks	None or Unspecified
Pad_Type	NaN
Ride_Control	None or Unspecified
Stick	NaN
Transmission	NaN
Turbocharged	NaN
Blade_Extension	NaN
Blade_Width	NaN
Enclosure_Type	NaN
Engine_Horsepower	NaN
Hydraulics	2 Valve
Pushblock	NaN
Ripper	NaN
Scarifier	NaN
Tip_Control	NaN
Tire_Size	None or Unspecified
Coupler	None or Unspecified
Coupler_System	NaN
Grouser_Tracks	NaN
Hydraulics_Flow	NaN
Track_Type	NaN
Undercarriage_Pad_Width	NaN
Stick_Length	NaN
Thumb	NaN
Pattern_Changer	NaN
Grouser_Type	NaN
Backhoe_Mounting	NaN
Blade_Type	NaN
Travel_Controls	NaN
Differential_Type	Standard
Steering_Controls	Conventional
saleYear	1989
saleMonth	1
saleDay	31
saleDayOfWeek	1
saleDayOfYear	31

		2 \
SalesID		1505138
SalePrice		50000.0
MachineID		1473654
ModelID		4139
datasource		132
auctioneerID		99.0
YearMade		1978
MachineHoursCurrentMeter		NaN
UsageBand		NaN
fiModelDesc		D7G
fiBaseModel		D7
fiSecondaryDesc		G
fiModelSeries		NaN
fiModelDescriptor		NaN
ProductSize		Large
fiProductClassDesc	Track Type Tractor, Dozer - 190.0 to 260.0 Hor...	
state		Florida
ProductGroup		TTT
ProductGroupDesc	Track Type Tractors	
Drive_System		NaN
Enclosure		OROPS
Forks		NaN
Pad_Type		NaN
Ride_Control		NaN
Stick		NaN
Transmission		Standard
Turbocharged		NaN
Blade_Extension		NaN
Blade_Width		NaN
Enclosure_Type		NaN
Engine_Horsepower		NaN
Hydraulics	2 Valve	
Pushblock		NaN
Ripper	None or Unspecified	
Scarifier		NaN
Tip_Control		NaN
Tire_Size		NaN
Coupler		NaN
Coupler_System		NaN
Grouser_Tracks		NaN
Hydraulics_Flow		NaN
Track_Type		NaN
Undercarriage_Pad_Width		NaN
Stick_Length		NaN
Thumb		NaN

Pattern_Changer		NaN
Grouser_Type		NaN
Backhoe_Mounting		None or Unspecified
Blade_Type		Straight
Travel_Controls		None or Unspecified
Differential_Type		NaN
Steering_Controls		NaN
saleYear		1989
saleMonth		1
saleDay		31
saleDayOfWeek		1
saleDayOfYear		31
	3	\
SalesID		1671174
SalePrice		16000.0
MachineID		1327630
ModelID		8591
datasource		132
auctioneerID		99.0
YearMade		1980
MachineHoursCurrentMeter		NaN
UsageBand		NaN
fiModelDesc		A62
fiBaseModel		A62
fiSecondaryDesc		NaN
fiModelSeries		NaN
fiModelDescriptor		NaN
ProductSize		NaN
fiProductClassDesc	Wheel Loader - Unidentified	
state		Florida
ProductGroup		WL
ProductGroupDesc		Wheel Loader
Drive_System		NaN
Enclosure		EROPS
Forks		None or Unspecified
Pad_Type		NaN
Ride_Control		None or Unspecified
Stick		NaN
Transmission		NaN
Turbocharged		NaN
Blade_Extension		NaN
Blade_Width		NaN
Enclosure_Type		NaN
Engine_Horsepower		NaN
Hydraulics	2 Valve	
Pushblock		NaN

Ripper		NaN
Scarifier		NaN
Tip_Control		NaN
Tire_Size	None or Unspecified	
Coupler	None or Unspecified	
Coupler_System		NaN
Grouser_Tracks		NaN
Hydraulics_Flow		NaN
Track_Type		NaN
Undercarriage_Pad_Width		NaN
Stick_Length		NaN
Thumb		NaN
Pattern_Changer		NaN
Grouser_Type		NaN
Backhoe_Mounting		NaN
Blade_Type		NaN
Travel_Controls		NaN
Differential_Type	Standard	
Steering_Controls	Conventional	
saleYear	1989	
saleMonth	1	
saleDay	31	
saleDayOfWeek	1	
saleDayOfYear	31	
		4
SalesID		1329056
SalePrice		22000.0
MachineID		1336053
ModelID		4089
datasource		132
auctioneerID		99.0
YearMade		1984
MachineHoursCurrentMeter		NaN
UsageBand		NaN
fiModelDesc		D3B
fiBaseModel		D3
fiSecondaryDesc		B
fiModelSeries		NaN
fiModelDescriptor		NaN
ProductSize		NaN
fiProductClassDesc	Track Type Tractor, Dozer - 20.0 to 75.0 Horse...	
state		Florida
ProductGroup		TTT
ProductGroupDesc	Track Type Tractors	
Drive_System		NaN
Enclosure		OROPS

Forks	NaN
Pad_Type	NaN
Ride_Control	NaN
Stick	NaN
Transmission	Standard
Turbocharged	NaN
Blade_Extension	NaN
Blade_Width	NaN
Enclosure_Type	NaN
Engine_Horsepower	NaN
Hydraulics	2 Valve
Pushblock	NaN
Ripper	None or Unspecified
Scarifier	NaN
Tip_Control	NaN
Tire_Size	NaN
Coupler	NaN
Coupler_System	NaN
Grouser_Tracks	NaN
Hydraulics_Flow	NaN
Track_Type	NaN
Undercarriage_Pad_Width	NaN
Stick_Length	NaN
Thumb	NaN
Pattern_Changer	NaN
Grouser_Type	NaN
Backhoe_Mounting	None or Unspecified
Blade_Type	PAT
Travel_Controls	Lever
Differential_Type	NaN
Steering_Controls	NaN
saleYear	1989
saleMonth	1
saleDay	31
saleDayOfWeek	1
saleDayOfYear	31

[171]: df_temp.isnull().sum()/len(df_temp)

SalesID	0.000000
SalePrice	0.000000
MachineID	0.000000
ModelID	0.000000
datasource	0.000000
auctioneerID	0.048791
YearMade	0.000000
MachineHoursCurrentMeter	0.642586

UsageBand	0.821492
fiModelDesc	0.000000
fiBaseModel	0.000000
fiSecondaryDesc	0.340993
fiModelSeries	0.857845
fiModelDescriptor	0.818715
ProductSize	0.524851
fiProductClassDesc	0.000000
state	0.000000
ProductGroup	0.000000
ProductGroupDesc	0.000000
Drive_System	0.740520
Enclosure	0.000809
Forks	0.520921
Pad_Type	0.803498
Ride_Control	0.629928
Stick	0.803498
Transmission	0.544444
Turbocharged	0.803498
Blade_Extension	0.937041
Blade_Width	0.937041
Enclosure_Type	0.937041
Engine_Horsepower	0.937041
Hydraulics	0.200062
Pushblock	0.937041
Ripper	0.740864
Scarifier	0.937014
Tip_Control	0.937041
Tire_Size	0.763415
Coupler	0.465277
Coupler_System	0.891024
Grouser_Tracks	0.891264
Hydraulics_Flow	0.891264
Track_Type	0.752378
Undercarriage_Pad_Width	0.750626
Stick_Length	0.752213
Thumb	0.752041
Pattern_Changer	0.752213
Grouser_Type	0.752378
Backhoe_Mounting	0.804428
Blade_Type	0.801610
Travel_Controls	0.801606
Differential_Type	0.826595
Steering_Controls	0.826697
saleYear	0.000000
saleMonth	0.000000
saleDay	0.000000

```
saleDayOfWeek          0.000000
saleDayOfYear          0.000000
dtype: float64
```

```
[172]: df_temp.isna().sum()
```

```
[172]: SalesID          0
SalePrice           0
MachineID          0
ModelID            0
datasource         0
auctioneerID       20136
YearMade           0
MachineHoursCurrentMeter 265194
UsageBand          339028
fiModelDesc        0
fiBaseModel        0
fiSecondaryDesc    140727
fiModelSeries      354031
fiModelDescriptor  337882
ProductSize        216605
fiProductClassDesc 0
state              0
ProductGroup       0
ProductGroupDesc   0
Drive_System       305611
Enclosure          334
Forks              214983
Pad_Type           331602
Ride_Control       259970
Stick              331602
Transmission       224691
Turbocharged       331602
Blade_Extension    386715
Blade_Width         386715
Enclosure_Type     386715
Engine_Horsepower  386715
Hydraulics          82565
Pushblock          386715
Ripper              305753
Scarifier           386704
Tip_Control        386715
Tire_Size           315060
Coupler             192019
Coupler_System     367724
Grouser_Tracks    367823
Hydraulics_Flow    367823
```

```

Track_Type           310505
Undercarriage_Pad_Width   309782
Stick_Length         310437
Thumb                310366
Pattern_Changer      310437
Grouser_Type         310505
Backhoe_Mounting     331986
Blade_Type           330823
Travel_Controls      330821
Differential_Type    341134
Steering_Controls    341176
saleYear              0
saleMonth             0
saleDay               0
saleDayOfWeek         0
saleDayOfYear         0
dtype: int64

```

1.7 Fill missing values

1.7.1 Fill numerical missing values first

[173]: df_temp.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 412698 entries, 0 to 412697
Data columns (total 57 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   SalesID           412698 non-null   int64  
 1   SalePrice          412698 non-null   float64 
 2   MachineID          412698 non-null   int64  
 3   ModelID            412698 non-null   int64  
 4   datasource          412698 non-null   int64  
 5   auctioneerID        392562 non-null   float64 
 6   YearMade            412698 non-null   int64  
 7   MachineHoursCurrentMeter 147504 non-null   float64 
 8   UsageBand           73670 non-null    object  
 9   fiModelDesc          412698 non-null   object  
 10  fiBaseModel          412698 non-null   object  
 11  fiSecondaryDesc       271971 non-null   object  
 12  fiModelSeries         58667 non-null    object  
 13  fiModelDescriptor      74816 non-null   object  
 14  ProductSize           196093 non-null   object  
 15  fiProductClassDesc      412698 non-null   object  
 16  state                 412698 non-null   object  
 17  ProductGroup          412698 non-null   object  
 18  ProductGroupDesc        412698 non-null   object 

```

```

19 Drive_System           107087 non-null  object
20 Enclosure              412364 non-null  object
21 Forks                  197715 non-null  object
22 Pad_Type                81096 non-null  object
23 Ride_Control            152728 non-null  object
24 Stick                   81096 non-null  object
25 Transmission             188007 non-null  object
26 Turbocharged            81096 non-null  object
27 Blade_Extension          25983 non-null  object
28 Blade_Width              25983 non-null  object
29 Enclosure_Type           25983 non-null  object
30 Engine_Horsepower        25983 non-null  object
31 Hydraulics              330133 non-null  object
32 Pushblock                25983 non-null  object
33 Ripper                  106945 non-null  object
34 Scarifier                25994 non-null  object
35 Tip_Control              25983 non-null  object
36 Tire_Size                 97638 non-null  object
37 Coupler                  220679 non-null  object
38 Coupler_System            44974 non-null  object
39 Grouser_Tracks           44875 non-null  object
40 Hydraulics_Flow           44875 non-null  object
41 Track_Type                102193 non-null  object
42 Undercarriage_Pad_Width    102916 non-null  object
43 Stick_Length               102261 non-null  object
44 Thumb                     102332 non-null  object
45 Pattern_Changer            102261 non-null  object
46 Grouser_Type               102193 non-null  object
47 Backhoe_Mounting            80712 non-null  object
48 Blade_Type                  81875 non-null  object
49 Travel_Controls              81877 non-null  object
50 Differential_Type            71564 non-null  object
51 Steering_Controls            71522 non-null  object
52 saleYear                  412698 non-null  int64
53 saleMonth                  412698 non-null  int64
54 saleDay                     412698 non-null  int64
55 saleDayOfWeek                412698 non-null  int64
56 saleDayOfYear                  412698 non-null  int64
dtypes: float64(3), int64(10), object(44)
memory usage: 179.5+ MB

```

```
[174]: for label,content in df_temp.items():
    if pd.api.types.is_object_dtype(content): # if read_csv we need to change
        ↪the dtype again to categorical
        df_temp[label] = content.astype("category").cat.as_ordered()
    if pd.api.types.is_string_dtype(content):
        df_temp[label] = content.astype("category").cat.as_ordered()
```

```

if pd.api.types.is_numeric_dtype(content):
    print(label)

# Hinweis: CSV-Dateien speichern keine Datentyp-Informationen (nur Werte als
# →Text).
# Beim pd.read_csv() werden die Spaltentypen deshalb von pandas neu "erraten"
# →(Type Inference),
# wodurch z.B. category, datetime oder Strings mit führenden Nullen als andere
# →dtypes eingelesen werden können.
# Lösung: dtypes/parse_dates beim Import explizit setzen oder ein Format wie
# →Parquet/Feather nutzen, das dtypes mit speichert.

```

SalesID
SalePrice
MachineID
ModelID
datasource
auctioneerID
YearMade
MachineHoursCurrentMeter
saleYear
saleMonth
saleDay
saleDayOfWeek
saleDayOfYear

[175]: df_temp.ModelID

```

[175]: 0      8434
       1      10150
       2      4139
       3      8591
       4      4089
       ...
      412693   5266
      412694   19330
      412695   17244
      412696   3357
      412697   4701
Name: ModelID, Length: 412698, dtype: int64

```

[176]: # Check for which numeric columns have null values

```

for label,content in df_temp.items():
    if pd.api.types.is_numeric_dtype(content):
        if pd.isnull(content).sum():
            print(label)

```

```

auctioneerID
MachineHoursCurrentMeter

[177]: # Fill numeric rows with the median

for label,content in df_temp.items():
    if pd.api.types.is_numeric_dtype(content):
        if pd.isnull(content).sum():
            #Add a binary column which tells us if the data was missing
            df_temp[label+"_is_missing"] = pd.isnull(content)
            #Fill missing numeric values with median
            df_temp[label]= content.fillna(content.median())

# Mean (arithm. Mittelwert) = Summe aller Werte / Anzahl: nutzt jede Beobachtung
# direkt und wird durch Ausreißer stark beeinflusst.
# Median = der mittlere Wert der sortierten Daten (bei gerader Anzahl: Mittel
# der zwei mittleren): robust gegenüber Ausreißern
# und oft besser für schiefe Verteilungen (z.B. Einkommen, Preise, Wartezeiten).

```

```
[178]: # Check for which numeric columns have null values
```

```

for label,content in df_temp.items():
    if pd.api.types.is_numeric_dtype(content):
        if pd.isnull(content).sum():
            print(label)

```

```
[179]: df_temp.auctioneerID_is_missing.value_counts()
```

```
[179]: auctioneerID_is_missing
False      392562
True       20136
Name: count, dtype: int64
```

```
[180]: df_temp.isna().sum()
```

SalesID	0
SalePrice	0
MachineID	0
ModelID	0
datasource	0
auctioneerID	0
YearMade	0
MachineHoursCurrentMeter	0
UsageBand	339028
fiModelDesc	0
fiBaseModel	0

fiSecondaryDesc	140727
fiModelSeries	354031
fiModelDescriptor	337882
ProductSize	216605
fiProductClassDesc	0
state	0
ProductGroup	0
ProductGroupDesc	0
Drive_System	305611
Enclosure	334
Forks	214983
Pad_Type	331602
Ride_Control	259970
Stick	331602
Transmission	224691
Turbocharged	331602
Blade_Extension	386715
Blade_Width	386715
Enclosure_Type	386715
Engine_Horsepower	386715
Hydraulics	82565
Pushblock	386715
Ripper	305753
Scarifier	386704
Tip_Control	386715
Tire_Size	315060
Coupler	192019
Coupler_System	367724
Grouser_Tracks	367823
Hydraulics_Flow	367823
Track_Type	310505
Undercarriage_Pad_Width	309782
Stick_Length	310437
Thumb	310366
Pattern_Changer	310437
Grouser_Type	310505
Backhoe_Mounting	331986
Blade_Type	330823
Travel_Controls	330821
Differential_Type	341134
Steering_Controls	341176
saleYear	0
saleMonth	0
saleDay	0
saleDayOfWeek	0
saleDayOfYear	0
auctioneerID_is_missing	0

```
MachineHoursCurrentMeter_is_missing          0
dtype: int64
```

```
[181]: ### Filling and turning categorical variables into numbers
```

```
#Check for columns which aren't numeric

for label,content in df_temp.items():
    if not pd.api.types.is_numeric_dtype(content):
        print(label)
```

```
UsageBand
fiModelDesc
fiBaseModel
fiSecondaryDesc
fiModelSeries
fiModelDescriptor
ProductSize
fiProductClassDesc
state
ProductGroup
ProductGroupDesc
Drive_System
Enclosure
Forks
Pad_Type
Ride_Control
Stick
Transmission
Turbocharged
Blade_Extension
Blade_Width
Enclosure_Type
Engine_Horsepower
Hydraulics
Pushblock
Ripper
Scarifier
Tip_Control
Tire_Size
Coupler
Coupler_System
Grouser_Tracks
Hydraulics_Flow
Track_Type
Undercarriage_Pad_Width
Stick_Length
```

```
Thumb
Pattern_Changer
Grouser_Type
Backhoe_Mounting
Blade_Type
Travel_Controls
Differential_Type
Steering_Controls
```

```
[182]: pd.Categorical(df_temp["state"]).dtype
```

```
[182]: CategoricalDtype(categories=['Alabama', 'Alaska', 'Arizona', 'Arkansas',
'California',
'Colorado', 'Connecticut', 'Delaware', 'Florida', 'Georgia',
'Hawaii', 'Idaho', 'Illinois', 'Indiana', 'Iowa', 'Kansas',
'Kentucky', 'Louisiana', 'Maine', 'Maryland',
'Massachusetts', 'Michigan', 'Minnesota', 'Mississippi',
'Missouri', 'Montana', 'Nebraska', 'Nevada', 'New Hampshire',
'New Jersey', 'New Mexico', 'New York', 'North Carolina',
'North Dakota', 'Ohio', 'Oklahoma', 'Oregon', 'Pennsylvania',
'Puerto Rico', 'Rhode Island', 'South Carolina',
'South Dakota', 'Tennessee', 'Texas', 'Unspecified', 'Utah',
'Vermont', 'Virginia', 'Washington', 'Washington DC',
'West Virginia', 'Wisconsin', 'Wyoming'],
, ordered=True, categories_dtype=object)
```

```
[183]: pd.Categorical(df_temp["state"]).codes # if missing value then -1
```

```
[183]: array([43,  8,  8, ...,  4,  4,  4], shape=(412698,), dtype=int8)
```

```
[184]: # Turn categorical variables into number and fill missing

for label,content in df_temp.items():
    if not pd.api.types.is_numeric_dtype(content):
        # add binary columns to indicate whether sample had missing value
        df_temp[label+"_is_missing"] = pd.isnull(content)
        #Turn categories into numbers and add +1
        df_temp[label] = pd.Categorical(content).codes+1
```

```
[185]: df_temp.head()
```

```
[185]:   SalesID  SalePrice  MachineID  ModelID  datasource  auctioneerID  YearMade \
0   1646770     9500.0    1126363      8434        132       18.0      1974
1   1821514    14000.0    1194089      10150        132       99.0      1980
2   1505138    50000.0    1473654      4139        132       99.0      1978
3   1671174    16000.0    1327630      8591        132       99.0      1980
4   1329056    22000.0    1336053      4089        132       99.0      1984
```

```

MachineHoursCurrentMeter  UsageBand  fiModelDesc ... \
0                      0.0        0      4593   ...
1                      0.0        0      1820   ...
2                      0.0        0      2348   ...
3                      0.0        0      1819   ...
4                      0.0        0      2119   ...

Undercarriage_Pad_Width_is_missing  Stick_Length_is_missing \
0                           True          True
1                           True          True
2                           True          True
3                           True          True
4                           True          True

Thumb_is_missing  Pattern_Changer_is_missing  Grouser_Type_is_missing \
0                 True          True          True
1                 True          True          True
2                 True          True          True
3                 True          True          True
4                 True          True          True

Backhoe_Mounting_is_missing  Blade_Type_is_missing \
0                   False         False
1                   True          True
2                   False         False
3                   True          True
4                   False         False

Travel_Controls_is_missing  Differential_Type_is_missing \
0                   False         True
1                   True          False
2                   False         True
3                   True          False
4                   False         True

Steering_Controls_is_missing
0                     True
1                     False
2                     True
3                     False
4                     True

[5 rows x 103 columns]

```

[186]: df_temp.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 412698 entries, 0 to 412697

```

```
Columns: 103 entries, SalesID to Steering_Controls_is_missing  
dtypes: bool(46), float64(3), int16(4), int64(10), int8(40)  
memory usage: 77.9 MB
```

```
[187]: df_temp.head().T
```

```
[187]:
```

	0	1	2	3	4
SalesID	1646770	1821514	1505138	1671174	1329056
SalePrice	9500.0	14000.0	50000.0	16000.0	22000.0
MachineID	1126363	1194089	1473654	1327630	1336053
ModelID	8434	10150	4139	8591	4089
datasource	132	132	132	132	132
...
Backhoe_Mounting_is_missing	False	True	False	True	False
Blade_Type_is_missing	False	True	False	True	False
Travel_Controls_is_missing	False	True	False	True	False
Differential_Type_is_missing	True	False	True	False	True
Steering_Controls_is_missing	True	False	True	False	True

[103 rows x 5 columns]

```
[188]: df_temp.isna().sum()[:70]
```

```
[188]:
```

SalesID	0
SalePrice	0
MachineID	0
ModelID	0
datasource	0
...	..
ProductSize_is_missing	0
fiProductClassDesc_is_missing	0
state_is_missing	0
ProductGroup_is_missing	0
ProductGroupDesc_is_missing	0

Length: 70, dtype: int64

Now that all of our data is numeric as well as our df has no missing values, we should be able to build a machine learning model

```
[189]: df_temp.head()
```

```
[189]:
```

	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	YearMade	\
0	1646770	9500.0	1126363	8434	132	18.0	1974	
1	1821514	14000.0	1194089	10150	132	99.0	1980	
2	1505138	50000.0	1473654	4139	132	99.0	1978	
3	1671174	16000.0	1327630	8591	132	99.0	1980	
4	1329056	22000.0	1336053	4089	132	99.0	1984	

```

MachineHoursCurrentMeter  UsageBand  fiModelDesc ... \
0                      0.0        0      4593   ...
1                      0.0        0      1820   ...
2                      0.0        0      2348   ...
3                      0.0        0      1819   ...
4                      0.0        0      2119   ...

Undercarriage_Pad_Width_is_missing  Stick_Length_is_missing \
0                           True          True
1                           True          True
2                           True          True
3                           True          True
4                           True          True

Thumb_is_missing  Pattern_Changer_is_missing  Grouser_Type_is_missing \
0                 True          True          True
1                 True          True          True
2                 True          True          True
3                 True          True          True
4                 True          True          True

Backhoe_Mounting_is_missing  Blade_Type_is_missing \
0                   False         False
1                   True          True
2                   False         False
3                   True          True
4                   False         False

Travel_Controls_is_missing  Differential_Type_is_missing \
0                   False         True
1                   True          False
2                   False         True
3                   True          False
4                   False         True

Steering_Controls_is_missing
0                     True
1                     False
2                     True
3                     False
4                     True

[5 rows x 103 columns]

```

[190]: %%time

#JN Function to calculate how much time the cell take

```
#Instantiate model
model = RandomForestRegressor(n_jobs = -1)

# Fit the model

model.fit(df_temp.drop("SalePrice", axis = 1), df_temp["SalePrice"])
```

CPU times: total: 32min 12s
Wall time: 5min 6s

[190]: RandomForestRegressor(n_jobs=-1)

[191]: # Score the model with train data wrong
model.score(df_temp.drop("SalePrice", axis = 1), df_temp["SalePrice"])

[191]: 0.9875665110400845

Question: Why doesn't the above metric reliable?

1.7.2 Splitting data into train/validation sets

[192]: df_temp.saleYear

```
[192]: 0          1989
      1          1989
      2          1989
      3          1989
      4          1989
      ...
412693     2012
412694     2012
412695     2012
412696     2012
412697     2012
Name: saleYear, Length: 412698, dtype: int64
```

[193]: df_temp.saleYear.value_counts()

```
[193]: saleYear
2009    43849
2008    39767
2011    35197
2010    33390
2007    32208
2006    21685
2005    20463
2004    19879
2001    17594
```

```

2000    17415
2002    17246
2003    15254
1998    13046
1999    12793
2012    11573
1997    9785
1996    8829
1995    8530
1994    7929
1993    6303
1992    5519
1991    5109
1989    4806
1990    4529
Name: count, dtype: int64

```

```

[194]: # 1) Split We have data leakage because we calculate the median with past and ↵ future datasets
df_val = df_temp[df_temp.saleYear == 2012].copy()
df_train = df_temp[df_temp.saleYear != 2012].copy()

# 2) Imputation-Statistiken nur aus Training
num_cols = df_train.select_dtypes(include="number").columns
medians = df_train[num_cols].median()

df_train[num_cols] = df_train[num_cols].fillna(medians)

# 2) Imputation-Statistiken nur aus Validation
num_cols = df_val.select_dtypes(include="number").columns
medians = df_val[num_cols].median()

df_val[num_cols] = df_val[num_cols].fillna(medians)

```

```

[195]: # Split data into X and y

X_train,y_train = df_train.drop("SalePrice",axis = 1), df_train.SalePrice

# Split data into X and y

X_valid,y_valid = df_val.drop("SalePrice",axis = 1), df_val.SalePrice

```

1.7.3 Build an evaluation function

```
[196]: # Create evaluation function (the competition uses RMSLE)
from sklearn.metrics import mean_squared_log_error, mean_absolute_error,r2_score

def rmsle (y_test,y_preds):
    """
    Calculates root mean squared log error between predictions and true labels.
    """
    return np.sqrt(mean_squared_log_error(y_test,y_preds))

# Create a function to evaluate model on a few diffrent levels
def show_scores(model):
    train_preds = model.predict(X_train)
    val_preds = model.predict(X_valid)
    scores = {"Training MAE":mean_absolute_error(y_train,train_preds),
              "Valid MAE": mean_absolute_error(y_valid,val_preds),
              "Training RMSLE":rmsle(y_train,train_preds),
              "Valid RMSLE":rmsle(y_valid,val_preds),
              "Training r^2":r2_score(y_train,train_preds),
              "Valid r^2":r2_score(y_valid,val_preds)}
    return scores
```

1.8 Testing our model on a subset (to tune the hyperparameters)

```
[197]: model = RandomForestRegressor(n_jobs = -1,
                                   random_state = 42)
```

```
#This takes far too long .... for experimenting
#model.fit(X_train,y_train)
```

```
[198]: len(X_train)
```

```
[198]: 401125
```

```
[199]: #Change max_samples value

model = RandomForestRegressor(n_jobs = -1, random_state = 42,max_samples=10000)
```

```
[200]: %%time
##cutting down on the max number of samples each estimators can see improves
    ↪training time
model.fit(X_train,y_train)
```

```
CPU times: total: 1min 5s
Wall time: 10.8 s
```

```
[200]: RandomForestRegressor(max_samples=10000, n_jobs=-1, random_state=42)
```

```
[201]: show_scores(model)
```

```
[201]: {'Training MAE': 5561.298809224058,
        'Valid MAE': 7177.26365505919,
        'Training RMSLE': np.float64(0.25774537825697696),
        'Valid RMSLE': np.float64(0.29362638671089003),
        'Training r^2': 0.8606658995199188,
        'Valid r^2': 0.8320374995090507}
```

1.8.1 Hyperparameter tuning with RandomizedSearchCV

```
[202]: %%time
from sklearn.model_selection import RandomizedSearchCV

#Diffrent RandomForestRegressor HP
rf_grid ={
    "n_estimators":np.arange(10,100,10),
    "max_depth": [None,3,5,10],
    "min_samples_split": np.arange(2,20,2),
    "min_samples_leaf":np.arange(1,20,2),
    "max_samples": [10000],
    "max_features": [0.5,1,"sqrt","auto"]
}

# Instantiate RandomizesSearchCV
rs_model = RandomizedSearchCV(RandomForestRegressor(n_jobs = -1, random_state = 42), param_distributions=rf_grid,
                               n_iter= 5, # very low because it take too long!
                               cv=5,
                               verbose = True)

#Fit the RSCV model
rs_model.fit(X_train,y_train)
```

Fitting 5 folds for each of 5 candidates, totalling 25 fits

CPU times: total: 1min 33s

Wall time: 36.5 s

```
[202]: RandomizedSearchCV(cv=5,
                          estimator=RandomForestRegressor(n_jobs=-1, random_state=42),
                          n_iter=5,
                          param_distributions={'max_depth': [None, 3, 5, 10],
                                               'max_features': [0.5, 1, 'sqrt',
                                                               'auto'],
                                               'max_samples': [10000],
                                               'min_samples_leaf': array([ 1,  3,  5,
                                              7,  9, 11, 13, 15, 17, 19]),
                                               'min_samples_split': array([ 2,  4,  6,
```

```
8, 10, 12, 14, 16, 18]),  
        'n_estimators': array([10, 20, 30, 40,  
50, 60, 70, 80, 90])},  
        verbose=True)
```

```
[203]: #find the best HP for the model  
  
rs_model.best_params_
```

```
[203]: {'n_estimators': np.int64(30),  
'min_samples_split': np.int64(10),  
'min_samples_leaf': np.int64(5),  
'max_samples': 10000,  
'max_features': 'sqrt',  
'max_depth': 10}
```

```
[204]: # Evaluate the randomized Search Model  
show_scores(rs_model)
```

```
[204]: {'Training MAE': 9024.775697210287,  
'Valid MAE': 10547.658931378968,  
'Training RMSLE': np.float64(0.3964490510077746),  
'Valid RMSLE': np.float64(0.41947442230936866),  
'Training r^2': 0.6780160387865647,  
'Valid r^2': 0.6808627887416161}
```

1.8.2 Train the model with the best hyperparameters

Note These were found after 100 iterations of RandomizedSearchCV.

```
[206]: %%time  
  
# Most ideal Hyperparameters  
  
ideal_model = RandomForestRegressor(n_estimators = 40,  
                                    min_samples_leaf=1,  
                                    min_samples_split=14,  
                                    max_features=0.5,  
                                    n_jobs = -1,  
                                    max_samples = None,  
                                    random_state = 42 # so our results are  
                           ↵reproducible  
)  
  
#Fit the ideal model  
ideal_model.fit(X_train,y_train)
```

```
CPU times: total: 5min 16s
Wall time: 53.6 s
```

```
[206]: RandomForestRegressor(max_features=0.5, min_samples_split=14, n_estimators=40,
                           n_jobs=-1, random_state=42)
```

```
[207]: show_scores(ideal_model)
```

```
[207]: {'Training MAE': 2953.816113716348,
         'Valid MAE': 5951.247761444453,
         'Training RMSLE': np.float64(0.14469006962371855),
         'Valid RMSLE': np.float64(0.2452416398953833),
         'Training r^2': 0.9588145522577225,
         'Valid r^2': 0.8818019502450093}
```

1.9 Make Predictions on test data

```
[208]: #Import the test data
```

```
df_test = pd.read_csv("data/Test.csv", low_memory = False,
                      parse_dates=["saledate"])
```

```
df_test.head()
```

```
[208]:   SalesID  MachineID  ModelID  datasource  auctioneerID  YearMade \
0    1227829      1006309      3168          121            3       1999
1    1227844      1022817      7271          121            3       1000
2    1227847      1031560     22805          121            3       2004
3    1227848       56204      1269          121            3       2006
4    1227863      1053887     22312          121            3       2005

   MachineHoursCurrentMeter  UsageBand  saledate  fiModelDesc  ... \
0             3688.0        Low  2012-05-03      580G  ...
1             28555.0       High  2012-05-10      936  ...
2             6038.0      Medium  2012-05-10  EC210BLC  ...
3             8940.0       High  2012-05-10     330CL  ...
4             2286.0        Low  2012-05-10      650K  ...

   Undercarriage_Pad_Width  Stick_Length  Thumb  Pattern_Changer \
0                NaN          NaN      NaN           NaN
1                NaN          NaN      NaN           NaN
2  None or Unspecified      9' 6"  Manual  None or Unspecified
3  None or Unspecified  None or Unspecified  Manual           Yes
4                NaN          NaN      NaN           NaN

   Grouser_Type  Backhoe_Mounting  Blade_Type  Travel_Controls \
0        NaN          NaN          NaN           NaN
1        NaN          NaN          NaN           NaN
```

```
2      Double          NaN          NaN          NaN
3      Triple          NaN          NaN          NaN
4      NaN  None or Unspecified      PAT  None or Unspecified
```

```
Differential_Type Steering_Controls
0            NaN          NaN
1      Standard    Conventional
2            NaN          NaN
3            NaN          NaN
4            NaN          NaN
```

[5 rows x 52 columns]

```
[209]: df_test.isna().sum()
```

```
SalesID                  0
MachineID                0
ModelID                  0
datasource                0
auctioneerID              0
YearMade                  0
MachineHoursCurrentMeter  10328
UsageBand                 10623
saledate                  0
fiModelDesc                0
fiBaseModel                0
fiSecondaryDesc            3975
fiModelSeries              10451
fiModelDescriptor           9433
ProductSize                6409
fiProductClassDesc          0
state                      0
ProductGroup                0
ProductGroupDesc             0
Drive_System                9698
Enclosure                  2
Forks                      6149
Pad_Type                   10349
Ride_Control                8216
Stick                      10349
Transmission                7639
Turbocharged                10349
Blade_Extension              11806
Blade_Width                  11806
Enclosure_Type                11806
Engine_Horsepower              11806
Hydraulics                  2142
```

```
Pushblock           11806
Ripper              9753
Scarifier           11806
Tip_Control         11806
Tire_Size            9679
Coupler              4856
Coupler_System       10391
Grouser_Tracks      10391
Hydraulics_Flow     10391
Track_Type           9063
Undercarriage_Pad_Width 9059
Stick_Length          9063
Thumb                 9062
Pattern_Changer       9063
Grouser_Type          9063
Backhoe_Mounting      10406
Blade_Type             10399
Travel_Controls        10399
Differential_Type      10328
Steering_Controls      10328
dtype: int64
```

```
[210]: df_test.dtypes
```

```
SalesID            int64
MachineID           int64
ModelID             int64
datasource           int64
auctioneerID         int64
YearMade             int64
MachineHoursCurrentMeter float64
UsageBand            object
saledate            datetime64[ns]
fiModelDesc          object
fiBaseModel          object
fiSecondaryDesc       object
fiModelSeries         object
fiModelDescriptor      object
ProductSize          object
fiProductClassDesc      object
state                object
ProductGroup          object
ProductGroupDesc       object
Drive_System           object
Enclosure             object
Forks                 object
Pad_Type               object
```

```
Ride_Control          object
Stick                object
Transmission         object
Turbocharged         object
Blade_Extension      object
Blade_Width          object
Enclosure_Type       object
Engine_Horsepower    object
Hydraulics           object
Pushblock             object
Ripper               object
Scarifier             object
Tip_Control          object
Tire_Size             object
Coupler               object
Coupler_System        object
Grouser_Tracks       object
Hydraulics_Flow      object
Track_Type            object
Undercarriage_Pad_Width object
Stick_Length          object
Thumb                object
Pattern_Changer      object
Grouser_Type          object
Backhoe_Mounting     object
Blade_Type            object
Travel_Controls       object
Differential_Type    object
Steering_Controls    object
dtype: object
```

```
[211]: df_test.columns
```

```
[211]: Index(['SalesID', 'MachineID', 'ModelID', 'datasource', 'auctioneerID',
       'YearMade', 'MachineHoursCurrentMeter', 'UsageBand', 'saledate',
       'fiModelDesc', 'fiBaseModel', 'fiSecondaryDesc', 'fiModelSeries',
       'fiModelDescriptor', 'ProductSize', 'fiProductClassDesc', 'state',
       'ProductGroup', 'ProductGroupDesc', 'Drive_System', 'Enclosure',
       'Forks', 'Pad_Type', 'Ride_Control', 'Stick', 'Transmission',
       'Turbocharged', 'Blade_Extension', 'Blade_Width', 'Enclosure_Type',
       'Engine_Horsepower', 'Hydraulics', 'Pushblock', 'Ripper', 'Scarifier',
       'Tip_Control', 'Tire_Size', 'Coupler', 'Coupler_System',
       'Grouser_Tracks', 'Hydraulics_Flow', 'Track_Type',
       'Undercarriage_Pad_Width', 'Stick_Length', 'Thumb', 'Pattern_Changer',
       'Grouser_Type', 'Backhoe_Mounting', 'Blade_Type', 'Travel_Controls',
       'Differential_Type', 'Steering_Controls'],
      dtype='object')
```

1.10 Preprocessing the data to the same format like X_Train

```
[212]: def preprocessing_data(df):
    """
    Performs transformation on df and returns transformed df.
    """
    df["saleYear"] = df.saledate.dt.year
    df["saleMonth"] = df.saledate.dt.month
    df["saleDay"] = df.saledate.dt.day
    df["saleDayOfWeek"] = df.saledate.dt.dayofweek
    df["saleDayOfYear"] = df.saledate.dt.dayofyear

    df.drop("saledate", axis = 1, inplace = True)

    # Fill the numeric rows with median
    for label,content in df.items():
        if pd.api.types.is_numeric_dtype(content):
            if pd.isnull(content).sum():
                #Add a binary column which tells us if the data was missing
                df[label+"_is_missing"] = pd.isnull(content)
                #Fill missing numeric values with median
                df[label]= content.fillna(content.median())

    #Filled categorical missing datat and turned categories into numbers

        if not pd.api.types.is_numeric_dtype(content):
            df[label+"_is_missing"] = pd.isnull(content)
            #We add +1 to the category code because pandas encodes missing
            #categories with -1
            df[label]=pd.Categorical(content).codes+1

    return df
```

```
[213]: df_test = preprocessing_data(df_test)
```

```
df_test.columns, df_test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12457 entries, 0 to 12456
Columns: 101 entries, SalesID to Steering_Controls_is_missing
dtypes: bool(45), float64(1), int16(2), int32(5), int64(6), int8(42)
memory usage: 2.0 MB
```

```
[213]: (Index(['SalesID', 'MachineID', 'ModelID', 'datasource', 'auctioneerID',
       'YearMade', 'MachineHoursCurrentMeter', 'UsageBand', 'fiModelDesc',
       'fiBaseModel',
       ...
       'Undercarriage_Pad_Width_is_missing', 'Stick_Length_is_missing',
```

```
'Thumb_is_missing', 'Pattern_Changer_is_missing',
'Grouser_Type_is_missing', 'Backhoe_Mounting_is_missing',
'Blade_Type_is_missing', 'Travel_Controls_is_missing',
'Differential_Type_is_missing', 'Steering_Controls_is_missing'],
dtype='object', length=101),
None)
```

[214]: # Make predictions on updated Test data dont work because a feature is missing
`#test_preds = ideal_model.predict(df_test)`

[215]: # We can find how the columns differ using sets
`set(X_train.columns)-set(df_test.columns)`

[215]: {'auctioneerID_is_missing'}

[216]: # Manually fit the column
`df_test["auctioneerID_is_missing"] = False`
`df_test.head()`

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	\
0	1227829	1006309	3168	121	3	1999	
1	1227844	1022817	7271	121	3	1000	
2	1227847	1031560	22805	121	3	2004	
3	1227848	56204	1269	121	3	2006	
4	1227863	1053887	22312	121	3	2005	

	MachineHoursCurrentMeter	UsageBand	fiModelDesc	fiBaseModel	...	\
0	3688.0	2	499	180	...	
1	28555.0	1	831	292	...	
2	6038.0	3	1177	404	...	
3	8940.0	1	287	113	...	
4	2286.0	2	566	196	...	

	Stick_Length_is_missing	Thumb_is_missing	Pattern_Changer_is_missing	\
0	True	True	True	
1	True	True	True	
2	False	False	False	
3	False	False	False	
4	True	True	True	

	Grouser_Type_is_missing	Backhoe_Mounting_is_missing	\
0	True	True	
1	True	True	
2	False	True	
3	False	True	
4	True	False	

```

Blade_Type_is_missing  Travel_Controls_is_missing  \
0                      True                      True
1                      True                      True
2                      True                      True
3                      True                      True
4                     False                     False

Differential_Type_is_missing  Steering_Controls_is_missing  \
0                           True                      True
1                          False                     False
2                           True                      True
3                           True                      True
4                           True                      True

auctioneerID_is_missing
0                      False
1                      False
2                      False
3                      False
4                      False

[5 rows x 102 columns]

```

```
[223]: # nach model.fit(X_train, y_train)
fit_cols = ideal_model.feature_names_in_

# X_test / X_new exakt wie beim Fit ausrichten
X_new_aligned = df_test.reindex(columns=fit_cols, fill_value=0)

test_preds = ideal_model.predict(X_new_aligned)
```

```
[222]: # We can find how the columns differ using sets
set(X_train.columns)-set(df_test.columns)
```

```
[222]: set()
```

```
[225]: len(test_preds)
```

```
[225]: 12457
```

```
[226]: test_preds
```

```
[226]: array([17030.00927386, 14355.53565165, 46623.08774286, ...,
           11964.85073347, 16496.71079281, 27119.99044029], shape=(12457,))
```

```
[227]: # Format predictions into the same format Kaggle is after
```

```
df_preds = pd.DataFrame()
df_preds ["SalesID"] = df_test["SalesID"]
df_preds ["SalesPrice"] = test_preds

df_preds
```

```
[227]:      SalesID      SalesPrice
```

```
0      1227829  17030.009274
1      1227844  14355.535652
2      1227847  46623.087743
3      1227848  71680.261335
4      1227863  61762.999424
...
12452   6643171  39966.363007
12453   6643173  12049.704433
12454   6643184  11964.850733
12455   6643186  16496.710793
12456   6643196  27119.990440
```

```
[12457 rows x 2 columns]
```

```
[228]: df_preds.to_csv("data/test_predictions.csv", index = False)
```

1.10.1 Feature Importance

seeks to figure out which different attributes of the data were most important when it comes to predicting the target variable

```
[231]: # Find feature importance of our best models
```

```
len(ideal_model.feature_importances_)
```

```
[231]: 102
```

```
[264]: # Helper function for plotting feature importance
```

```
def plot_features(columns, importanceFactor, top_n=20):
    columns = np.array(columns)
    importanceFactor = np.array(importanceFactor)

    # sortiere nach Wichtigkeit (absteigend) und nimm Top-N
    idx = np.argsort(importanceFactor)[::-1][:top_n]
    cols = columns[idx]
    imps = importanceFactor[idx]

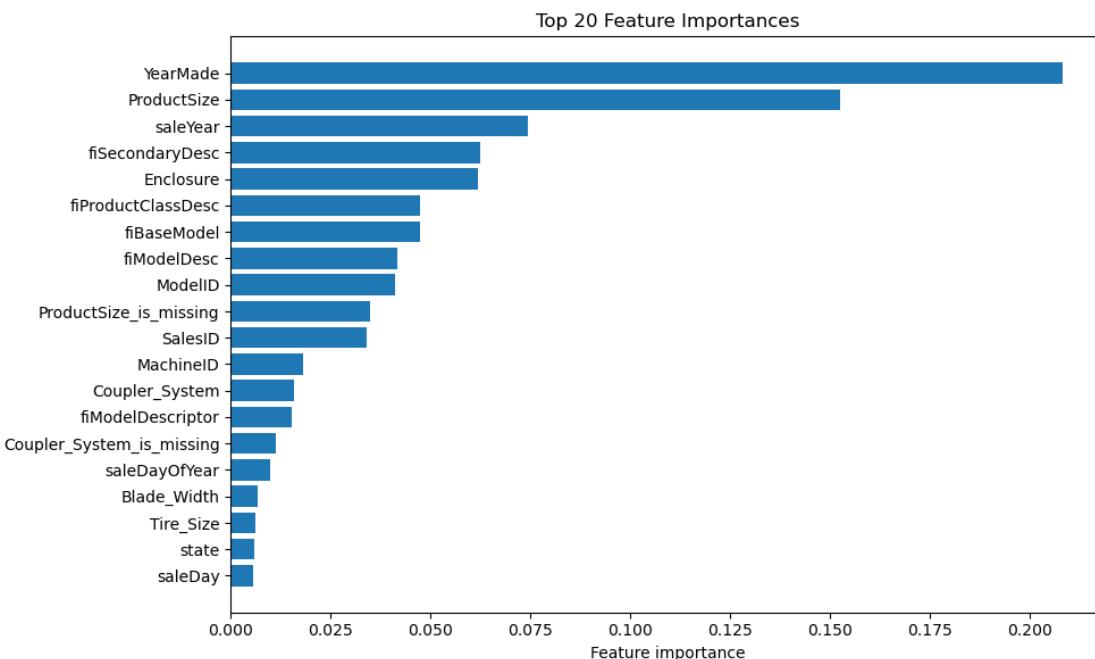
    fig, ax = plt.subplots(figsize=(10, 6))
    ax.barh(cols[::-1], imps[::-1]) # barh + umdrehen => wichtigste oben
```

```

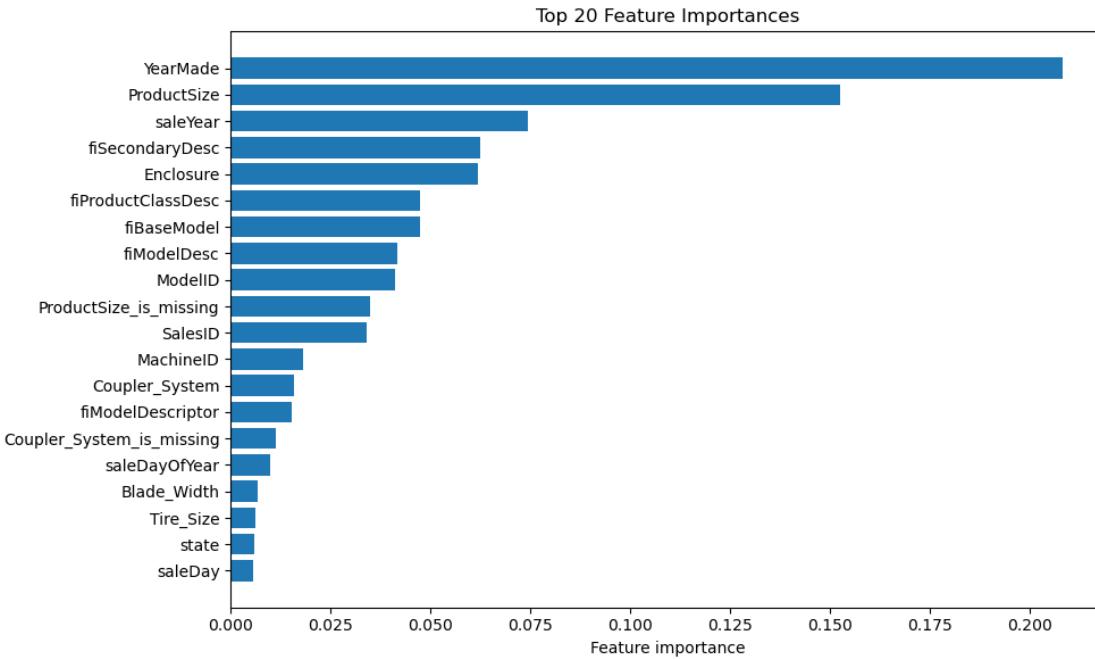
ax.set_xlabel("Feature importance")
ax.set_title(f"Top {top_n} Feature Importances")
plt.tight_layout()
plt.show()

plot_features(X_new_aligned.columns, ideal_model.feature_importances_, top_n=20)

```



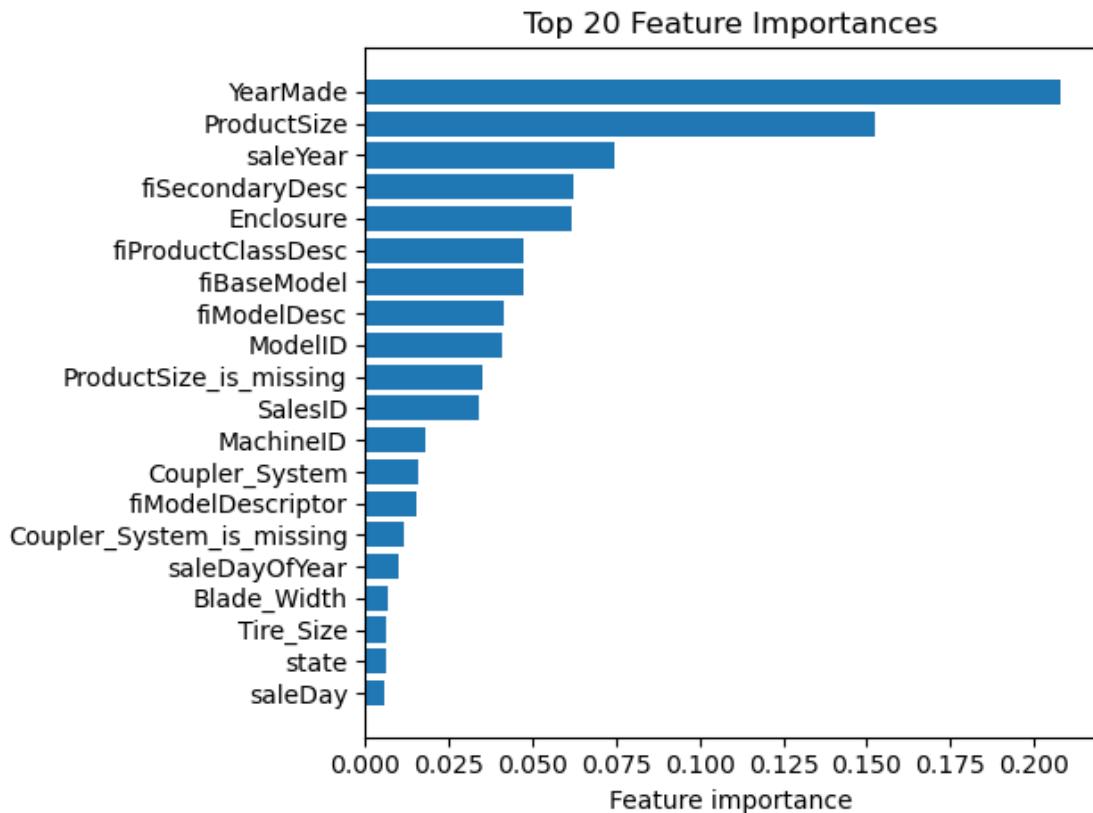
[265]: plot_features(X_new_aligned.columns,ideal_model.feature_importances_)



```
[277]: def plot_feature_other_way (columns, importances, n = 20):
    df = (pd.DataFrame({"features":columns,
                        "features_importances": importances})
          .sort_values("features_importances", ascending = False)
          .reset_index(drop=True))

    # Plot the dataframe
    fig, ax = plt.subplots()
    #ax.barh(np.flip(np.array(df["features"][:n:])), np.flip(np.
    ↵array(df["features_importances"][:n:])))
    ax.barh(df["features"][:n:][::-1], df["features_importances"][:n:][::-1])
    ax.set_xlabel("Feature importance")
    ax.set_title(f"Top {n} Feature Importances")
    plt.tight_layout()
    plt.show()
```

```
[278]: plot_feature_other_way(X_new_aligned.columns,ideal_model.feature_importances_)
```



Question to finish: Why might knowing the feature importances of a trained machine learning model be helpful?

Final challenge: What other machine learning models could you try on our dataset? Hint checkout the regression section of scikit learn map or try to look at CatBoost.ai or XGBoost.ai

[]: