

end-to-end-heart-disease-classification

February 11, 2026

1 Predicting heart disease using machine learning

This notebook looks into using various Python based machine learning and data science libraries in an attempt to build a machine learning model capable of predicting whether or not someone has heart disease based on their medical attributes

We're going to take the following approach: 1. Problem definition 2. Data 3. Evaluation 4. Features 5. Modelling 6. Experimentation

1.1 1. Problem Definition

In a statement, > Given clinical parameters about a patient, can we predict whether or not the have heart disease?

1.2 2. Data

The original data come from: <https://www.kaggle.com/datasets/redwankarimsony/heart-disease-data>

1.3 3. Evaluation

If we can reach 95% accuracy at predicting whether or not a patient has heart disease during the proof of concept we'll pursue the project

1.4 4. Features

This is where you will get diffrent information about each of the features of your data.

Create Data Dictionary

1. id (Unique id for each patient)
2. age (Age of the patient in years)
3. origin (place of study)
4. sex (Male/Female)
5. cp chest pain type ([typical angina, atypical angina, non-anginal, asymptomatic])
6. trestbps resting blood pressure (resting blood pressure (in mm Hg on admission to the hospital))
7. chol (serum cholesterol in mg/dl)
8. fbs (if fasting blood sugar > 120 mg/dl)
9. restecg (resting electrocardiographic results)
-- Values: [normal, stt abnormality, lv hypertrophy]
10. thalach: maximum heart rate achieved
11. exang: exercise-induced angina (True/ False)

12. oldpeak: ST depression induced by exercise relative to rest
13. slope: the slope of the peak exercise ST segment
14. ca: number of major vessels (0-3) colored by fluoroscopy
15. thal: [normal; fixed defect; reversible defect]
16. target: the predicted attribute

1.5 Preparing the tools

We are going to use pandas, matplotlib and NumPy for data analysis and manipulation

```
[72]: # Import all the tools

# Regular EDA (exploring data analysis) and plotting libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Models from scikit learn
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier

# Model Evaluations
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.model_selection import RandomizedSearchCV, GridSearchCV
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.metrics import precision_score, f1_score, recall_score
from sklearn.metrics import RocCurveDisplay
```

1.6 Load Data

```
[2]: df = pd.read_csv("./data/heart-disease.csv")
df.shape # rows, columns
```

```
[2]: (303, 14)
```

```
[3]: df.head()
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	63	1	3	145	233	1	0	150	0	2.3	0	
1	37	1	2	130	250	0	1	187	0	3.5	0	
2	41	0	1	130	204	0	0	172	0	1.4	2	
3	56	1	1	120	236	0	1	178	0	0.8	2	
4	57	0	0	120	354	0	1	163	1	0.6	2	

	ca	thal	target
0	0	1	1

```
1   0    2    1
2   0    2    1
3   0    2    1
4   0    2    1
```

1.7 Data Exploration (exploratory data analysis or EDA)

The goal here is to find out more about the data and become a subject matter expert on the dataset you are working with.

1. What questions are you trying to solve?
2. What kind of data do we have and how do we treat different types?
3. What's missing from the data and how do you deal with it
4. Where are the outliers and why should you care about them?
5. How can you add, change or remove features to get more out of your data

```
[4]: df.tail()
```

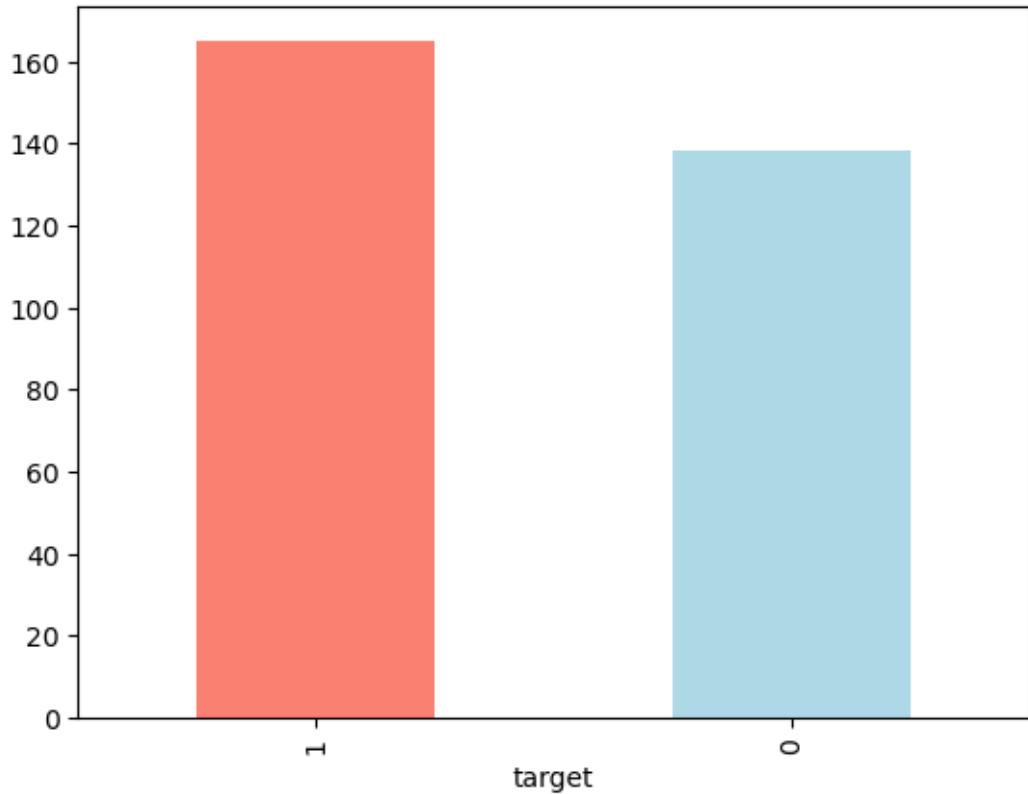
```
[4]:      age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak \
298    57    0    0     140    241    0        1     123      1      0.2
299    45    1    3     110    264    0        1     132      0      1.2
300    68    1    0     144    193    1        1     141      0      3.4
301    57    1    0     130    131    0        1     115      1      1.2
302    57    0    1     130    236    0        0     174      0      0.0

      slope  ca  thal  target
298      1    0    3      0
299      1    0    3      0
300      1    2    3      0
301      1    1    3      0
302      1    1    2      0
```

```
[5]: # Let's find out how many of each class there
df["target"].value_counts()
```

```
[5]: target
1    165
0    138
Name: count, dtype: int64
```

```
[6]: df["target"].value_counts().plot(kind = "bar", color=["salmon","lightblue"]);
```



[7]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
 ---  --          -----          ----- 
 0   age         303 non-null    int64  
 1   sex         303 non-null    int64  
 2   cp          303 non-null    int64  
 3   trestbps   303 non-null    int64  
 4   chol        303 non-null    int64  
 5   fbs         303 non-null    int64  
 6   restecg    303 non-null    int64  
 7   thalach    303 non-null    int64  
 8   exang       303 non-null    int64  
 9   oldpeak    303 non-null    float64 
 10  slope       303 non-null    int64  
 11  ca          303 non-null    int64  
 12  thal        303 non-null    int64  
 13  target      303 non-null    int64  
dtypes: float64(1), int64(13)
```

```
memory usage: 33.3 KB
```

```
[8]: #Are there any missing values ?  
df.isna().sum()
```

```
[8]: age      0  
sex       0  
cp        0  
trestbps  0  
chol      0  
fbs       0  
restecg   0  
thalach   0  
exang     0  
oldpeak   0  
slope     0  
ca        0  
thal      0  
target    0  
dtype: int64
```

```
[9]: df.describe()
```

```
[9]:          age      sex      cp      trestbps      chol      fbs \\\ncount  303.000000  303.000000  303.000000  303.000000  303.000000  303.000000  
mean    54.366337  0.683168  0.966997  131.623762  246.264026  0.148515  
std     9.082101  0.466011  1.032052  17.538143  51.830751  0.356198  
min    29.000000  0.000000  0.000000  94.000000  126.000000  0.000000  
25%   47.500000  0.000000  0.000000  120.000000  211.000000  0.000000  
50%   55.000000  1.000000  1.000000  130.000000  240.000000  0.000000  
75%   61.000000  1.000000  2.000000  140.000000  274.500000  0.000000  
max   77.000000  1.000000  3.000000  200.000000  564.000000  1.000000  
  
          restecg      thalach      exang      oldpeak      slope      ca \\\ncount  303.000000  303.000000  303.000000  303.000000  303.000000  303.000000  
mean    0.528053  149.646865  0.326733  1.039604  1.399340  0.729373  
std     0.525860  22.905161  0.469794  1.161075  0.616226  1.022606  
min    0.000000  71.000000  0.000000  0.000000  0.000000  0.000000  
25%   0.000000  133.500000  0.000000  0.000000  1.000000  0.000000  
50%   1.000000  153.000000  0.000000  0.800000  1.000000  0.000000  
75%   1.000000  166.000000  1.000000  1.600000  2.000000  1.000000  
max   2.000000  202.000000  1.000000  6.200000  2.000000  4.000000  
  
          thal      target  
count  303.000000  303.000000  
mean    2.313531  0.544554  
std     0.612277  0.498835  
min    0.000000  0.000000
```

```
25%      2.000000  0.000000
50%      2.000000  1.000000
75%      3.000000  1.000000
max      3.000000  1.000000
```

1.7.1 Heart Disease Frequency according to Sex

```
[10]: df.sex.value_counts()
```

```
[10]: sex
1    207
0     96
Name: count, dtype: int64
```

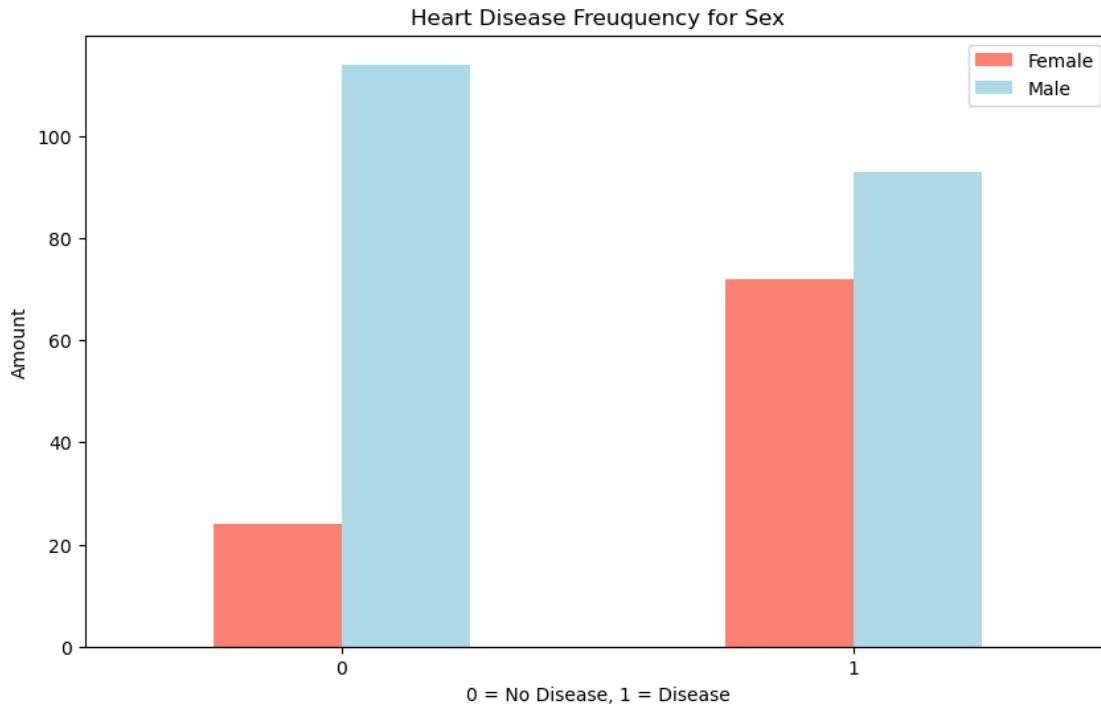
```
[11]: # Compare target column with sex column
pd.crosstab(df.target, df.sex)
```

```
[11]: sex      0      1
target
0        24   114
1        72   93
```

```
[12]: # Create a plot of crosstab
pd.crosstab(df.target, df.sex).plot(kind = "bar", figsize=(10,6),
                                     color=["salmon","lightblue"])

plt.title("Heart Disease Frequency for Sex")
plt.xlabel("0 = No Disease, 1 = Disease")
plt.ylabel("Amount")
plt.legend(["Female", "Male"])

plt.xticks(rotation=0);
```



```
[13]: df.head()
```

```
[13]:    age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  slope \
0    63    1    3      145   233    1        0     150      0       2.3      0
1    37    1    2      130   250    0        1     187      0       3.5      0
2    41    0    1      130   204    0        0     172      0       1.4      2
3    56    1    1      120   236    0        1     178      0       0.8      2
4    57    0    0      120   354    0        1     163      1       0.6      2

      ca  thal  target
0    0    1      1
1    0    2      1
2    0    2      1
3    0    2      1
4    0    2      1
```

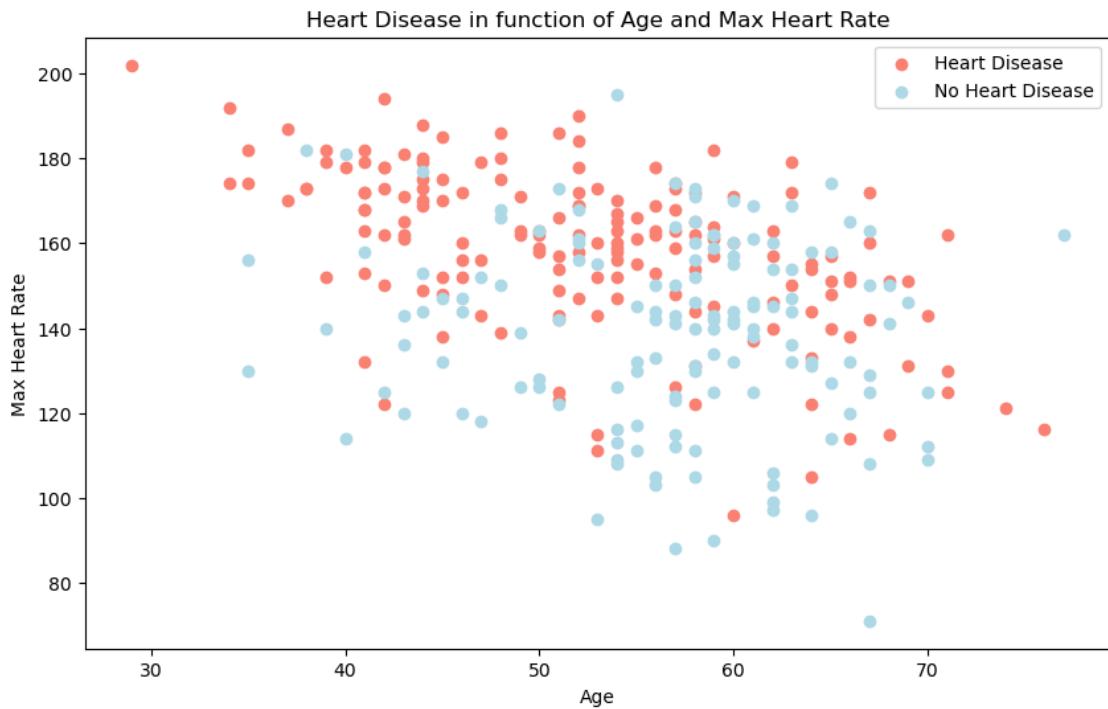
```
[14]: df["thalach"].value_counts()
```

```
[14]: thalach
162    11
160     9
163     9
152     8
173     8
```

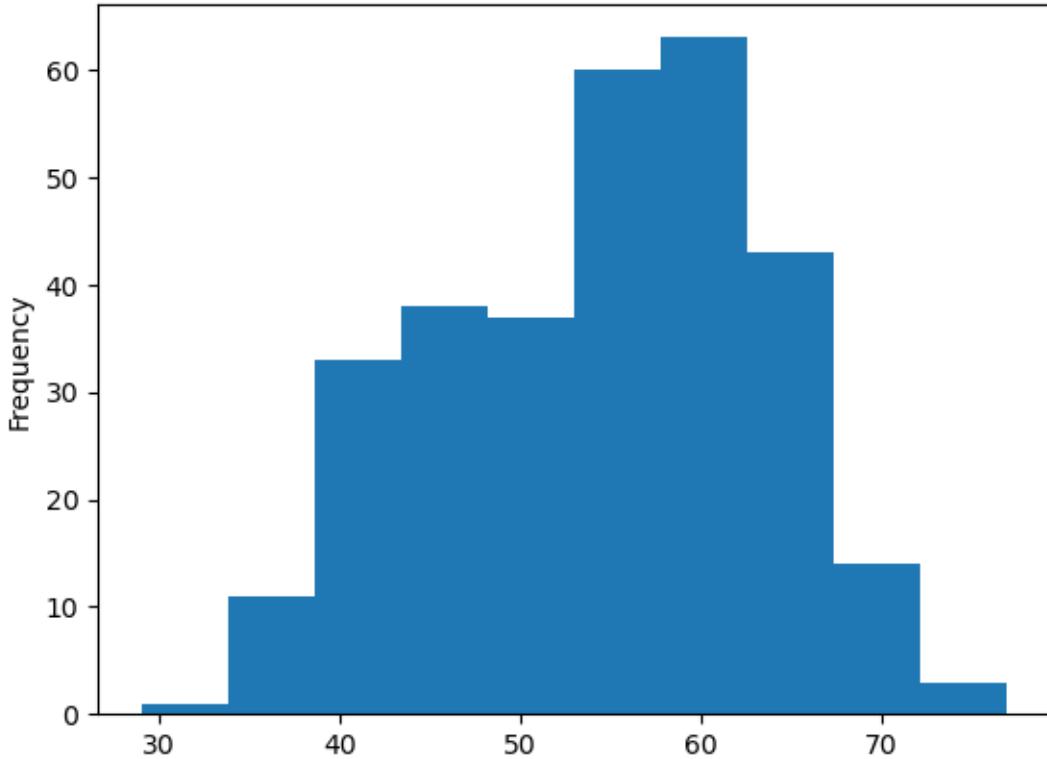
```
..  
202      1  
184      1  
121      1  
192      1  
90       1  
Name: count, Length: 91, dtype: int64
```

1.8 Age vs Max Heart Rate or Heart Disease

```
[15]: # Create another figure  
plt.figure(figsize= (10,6))  
  
# Scatter with positive examples  
plt.scatter(df.age[df.target ==1],  
            df.thalach[df.target == 1],  
            c = "salmon");  
  
# Scatter with negative examples  
plt.scatter(df.age[df.target ==0],  
            df.thalach[df.target == 0],  
            c = "lightblue");  
  
plt.legend(["Heart Disease","No Heart Disease"]);  
plt.title("Heart Disease in function of Age and Max Heart Rate")  
plt.xlabel("Age")  
plt.ylabel("Max Heart Rate");
```



```
[16]: # Check the distribution of the age column with a histogram  
df.age.plot.hist();
```



1.9 Heart Disease Frequency per Chest Pain Type

- 5. cp chest pain type ([typical angina, atypical angina, non-anginal, asymptomatic])

```
[17]: pd.crosstab(df.cp,df.target)
```

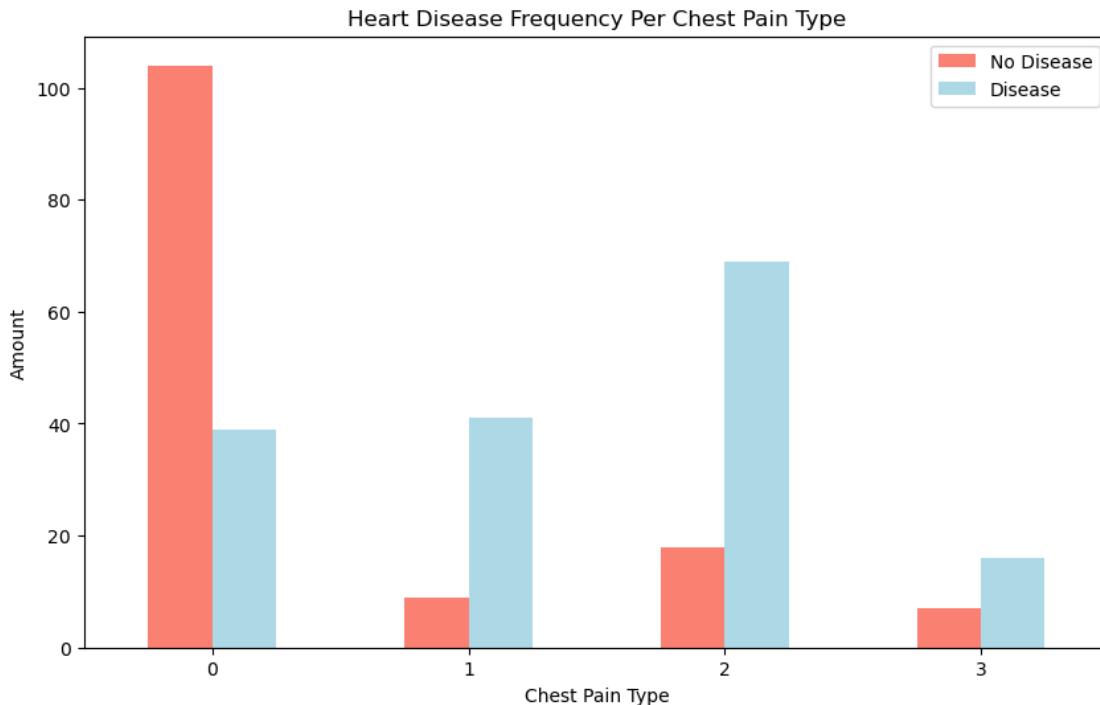
```
[17]: target      0      1
cp
0          104    39
1           9    41
2          18    69
3           7    16
```

```
[18]: # Make the crosstab more visual

pd.crosstab(df.cp,df.target).plot(kind="bar",
                                    figsize=(10,6),
                                    color= ["salmon", "lightblue"])

# Add some communication
plt.title("Heart Disease Frequency Per Chest Pain Type")
plt.xlabel("Chest Pain Type")
```

```
plt.ylabel("Amount")
plt.legend(["No Disease", "Disease"])
plt.xticks(rotation=0);
```



```
[19]: df.head()
```

```
[19]:    age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  slope \
0    63    1    3      145   233    1      0       150      0     2.3      0
1    37    1    2      130   250    0      1       187      0     3.5      0
2    41    0    1      130   204    0      0       172      0     1.4      2
3    56    1    1      120   236    0      1       178      0     0.8      2
4    57    0    0      120   354    0      1       163      1     0.6      2

      ca  thal  target
0    0    1      1
1    0    2      1
2    0    2      1
3    0    2      1
4    0    2      1
```

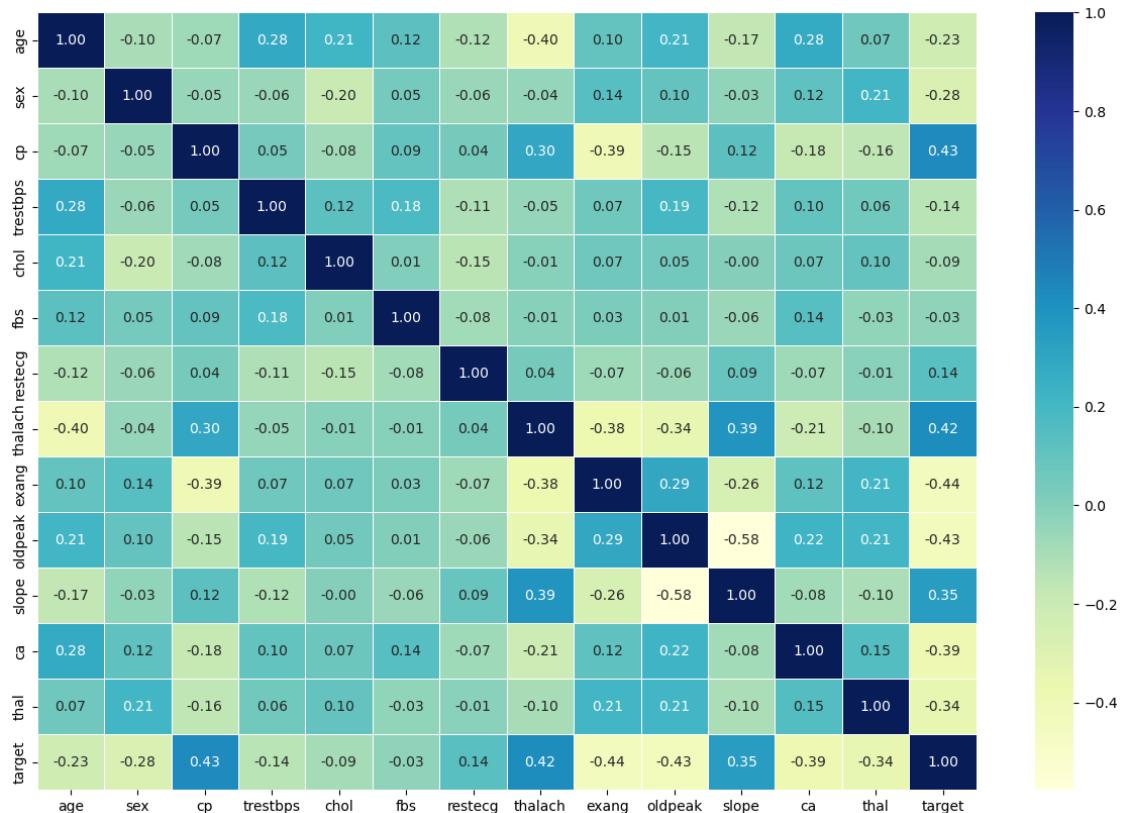
```
[20]: # Make a correlation matrix
df.corr()
```

[20] :

	age	sex	cp	trestbps	chol	fbs	\
age	1.000000	-0.098447	-0.068653	0.279351	0.213678	0.121308	
sex	-0.098447	1.000000	-0.049353	-0.056769	-0.197912	0.045032	
cp	-0.068653	-0.049353	1.000000	0.047608	-0.076904	0.094444	
trestbps	0.279351	-0.056769	0.047608	1.000000	0.123174	0.177531	
chol	0.213678	-0.197912	-0.076904	0.123174	1.000000	0.013294	
fbs	0.121308	0.045032	0.094444	0.177531	0.013294	1.000000	
restecg	-0.116211	-0.058196	0.044421	-0.114103	-0.151040	-0.084189	
thalach	-0.398522	-0.044020	0.295762	-0.046698	-0.009940	-0.008567	
exang	0.096801	0.141664	-0.394280	0.067616	0.067023	0.025665	
oldpeak	0.210013	0.096093	-0.149230	0.193216	0.053952	0.005747	
slope	-0.168814	-0.030711	0.119717	-0.121475	-0.004038	-0.059894	
ca	0.276326	0.118261	-0.181053	0.101389	0.070511	0.137979	
thal	0.068001	0.210041	-0.161736	0.062210	0.098803	-0.032019	
target	-0.225439	-0.280937	0.433798	-0.144931	-0.085239	-0.028046	
	restecg	thalach	exang	oldpeak	slope	ca	\
age	-0.116211	-0.398522	0.096801	0.210013	-0.168814	0.276326	
sex	-0.058196	-0.044020	0.141664	0.096093	-0.030711	0.118261	
cp	0.044421	0.295762	-0.394280	-0.149230	0.119717	-0.181053	
trestbps	-0.114103	-0.046698	0.067616	0.193216	-0.121475	0.101389	
chol	-0.151040	-0.009940	0.067023	0.053952	-0.004038	0.070511	
fbs	-0.084189	-0.008567	0.025665	0.005747	-0.059894	0.137979	
restecg	1.000000	0.044123	-0.070733	-0.058770	0.093045	-0.072042	
thalach	0.044123	1.000000	-0.378812	-0.344187	0.386784	-0.213177	
exang	-0.070733	-0.378812	1.000000	0.288223	-0.257748	0.115739	
oldpeak	-0.058770	-0.344187	0.288223	1.000000	-0.577537	0.222682	
slope	0.093045	0.386784	-0.257748	-0.577537	1.000000	-0.080155	
ca	-0.072042	-0.213177	0.115739	0.222682	-0.080155	1.000000	
thal	-0.011981	-0.096439	0.206754	0.210244	-0.104764	0.151832	
target	0.137230	0.421741	-0.436757	-0.430696	0.345877	-0.391724	
	thal	target					
age	0.068001	-0.225439					
sex	0.210041	-0.280937					
cp	-0.161736	0.433798					
trestbps	0.062210	-0.144931					
chol	0.098803	-0.085239					
fbs	-0.032019	-0.028046					
restecg	-0.011981	0.137230					
thalach	-0.096439	0.421741					
exang	0.206754	-0.436757					
oldpeak	0.210244	-0.430696					
slope	-0.104764	0.345877					
ca	0.151832	-0.391724					
thal	1.000000	-0.344029					
target	-0.344029	1.000000					

```
[21]: # Let's make the correlation matrix a little bit prettier
```

```
corr_matrix = df.corr()
fig,ax = plt.subplots(figsize=(15,10))
ax = sns.heatmap(corr_matrix,
                  annot=True,
                  linewidths=0.5,
                  fmt=".2f",
                  cmap="YlGnBu");
```



2 5. Modeling

```
[22]: from sklearn.pipeline import Pipeline
```

```
# Split the data
X = df.drop("target",axis = 1)
y = df["target"]

#reproduce the results
np.random.seed(42)
```

```
X_train,X_test,y_train,y_test = train_test_split(X,y,shuffle=True, test_size = 0.  
→2)
```

[24]: X_train.shape, X_test.shape

[24]: ((242, 13), (61, 13))

Now we have got our datat split into training and test sets, it's time to build a machine learning model. We will train it (find the patterns) on the training set. And we will test it (use the patterns) on the test set.

We are going to try 3 diffrent machine learning models: 1. Logistic Regression 2. K-Nearest Neighbours Classifier 3. Random Forest Classifier

[25]: #Put models into a Dict

```
models = { "Logistic Regression": LogisticRegression(),  
          "KNN": KNeighborsClassifier(),  
          "Random Forest": RandomForestClassifier()}  
  
# Create a function to fit and score our models  
def fit_and_score (models, X_train,X_test,y_train,y_test):  
    """  
    Fits and evaluates given machine learning models.  
    models: a dict of diffrent Scikit.Learn machine learning models  
    X_train: training data (no labels)  
    X_test: test data (no labels)  
    y_train: training data (labels)  
    y_test: test data (labels)  
  
    """  
    np.random.seed(42)  
  
    # Make a dict to keep model scores  
  
    model_scores={}  
  
    # Loop through models  
    for name, model in models.items():  
        #Fit the model to the data  
        model.fit(X_train,y_train)  
        #Evaluate the model and append it to the model_scores dict  
        model_scores[name] = model.score(X_test,y_test)  
    return model_scores
```

[26]: models_scores = fit_and_score(models,X_train,X_test,y_train,y_test)

C:\Users\RomaM\anaconda3\Lib\site-

```
packages\sklearn\linear_model\_logistic.py:473: ConvergenceWarning: lbfgs failed  
to converge after 100 iteration(s) (status=1):  
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT
```

Increase the number of iterations to improve the convergence (max_iter=100).

You might also want to scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

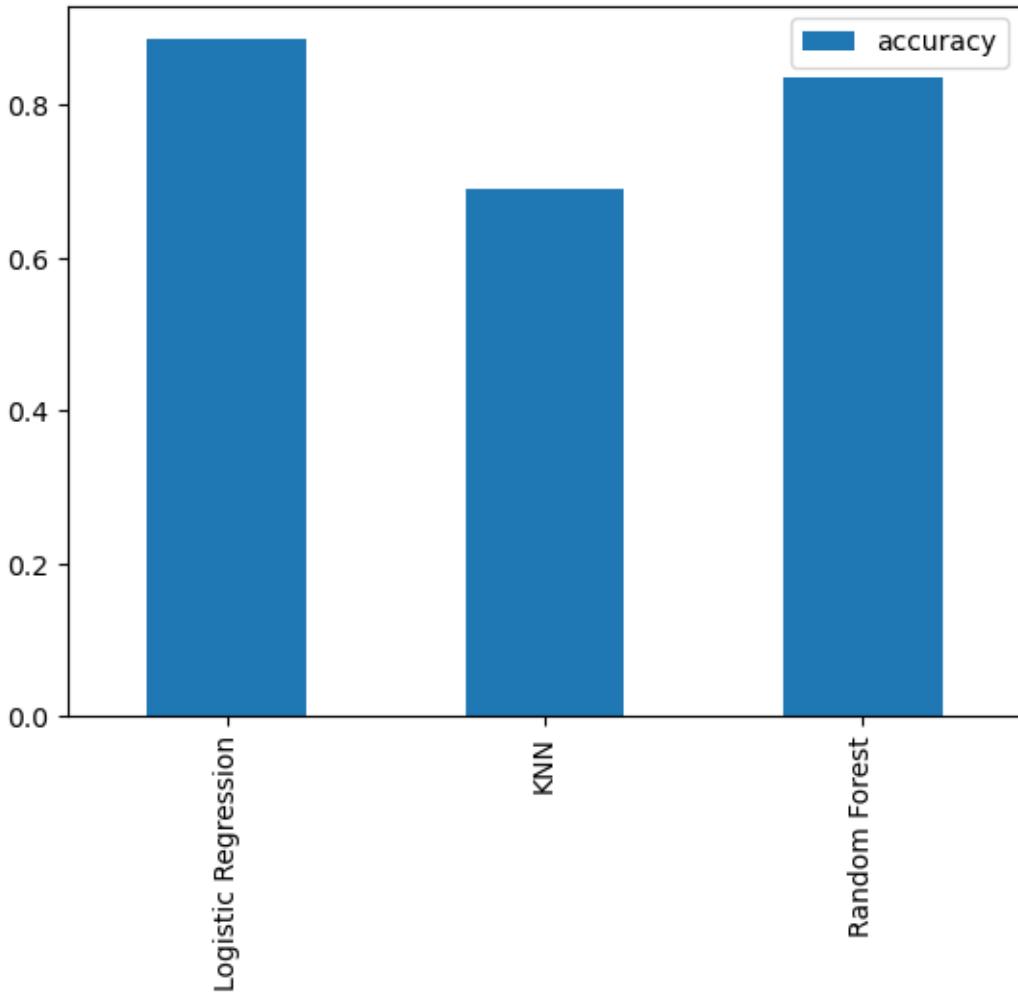
```
n_iter_i = _check_optimize_result()
```

[27]: models_scores

```
[27]: {'Logistic Regression': 0.8852459016393442,  
'KNN': 0.6885245901639344,  
'Random Forest': 0.8360655737704918}
```

2.1 Model Comparison

[28]: model_compare = pd.DataFrame(models_scores, index = ["accuracy"])
model_compare.T.plot.bar();



```
[29]: model_compare
```

```
[29]:          Logistic Regression      KNN  Random Forest
accuracy      0.885246   0.688525      0.836066
```

Now we have got a baseline model and we know a model's first predictions aren't always what we should base our next steps off. What should we do?

Let's look at the following:

- Hyperparameter tuning
- Feature importance
- Confusion Matrix
- Cross_Validation
- Precision
- Recall
- F1 score

- Classification report
- ROC curve
- Area under the curve (AUC)

2.2 Hyperparameter Tuning (by hand)

[30]: # Let's tune KNN

```
train_scores = []
test_scores = []

# Create a list of diffrent values for n neighbors
neighbors = range (1,21)

# Setup KNN instance
knn = KNeighborsClassifier();

#Loop through diffrent n_neigbors

for i in neighbors:
    knn.set_params(n_neighbors=i)

    # Fit the alg
    knn.fit(X_train,y_train)

    # Updating the training scores list
    train_scores.append(knn.score(X_train,y_train))

    # Update the test scores list
    test_scores.append(knn.score(X_test,y_test))
```

[31]: train_scores

```
[31]: [1.0,
 0.8099173553719008,
 0.7727272727272727,
 0.743801652892562,
 0.7603305785123967,
 0.7520661157024794,
 0.743801652892562,
 0.7231404958677686,
 0.71900826446281,
 0.6942148760330579,
 0.7272727272727273,
 0.6983471074380165,
 0.6900826446280992,
```

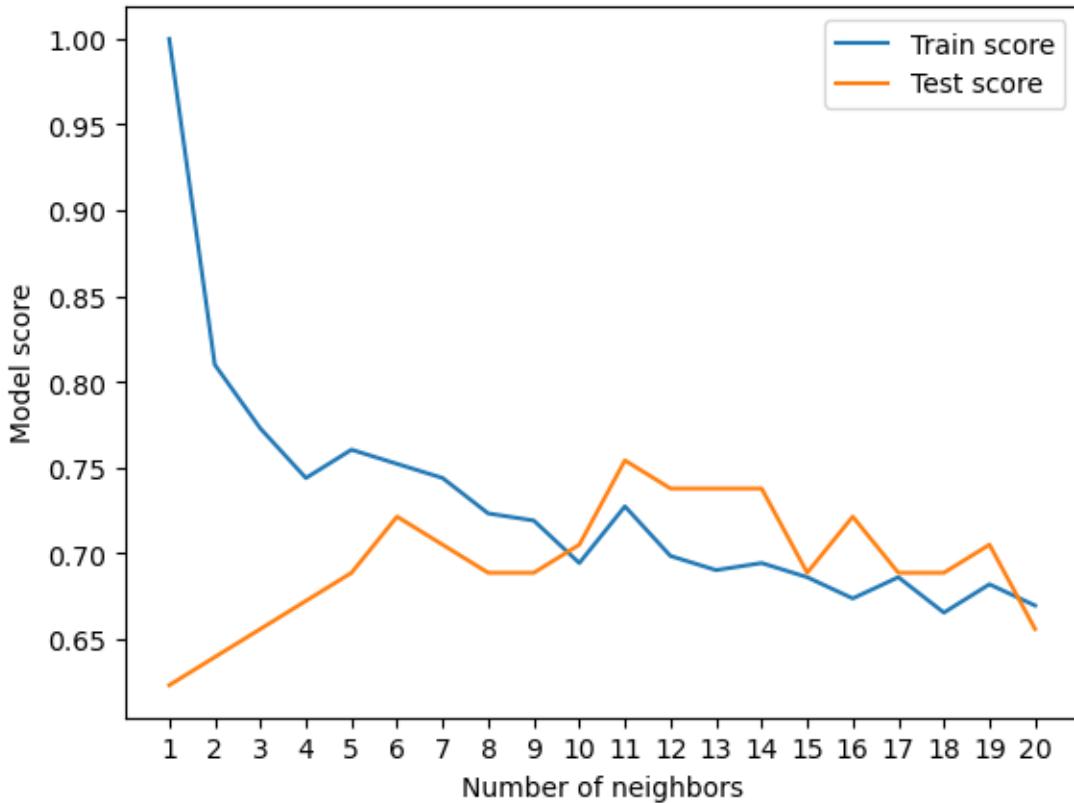
```
0.6942148760330579,  
0.6859504132231405,  
0.6735537190082644,  
0.6859504132231405,  
0.6652892561983471,  
0.6818181818181818,  
0.6694214876033058]
```

```
[32]: test_scores
```

```
[32]: [0.6229508196721312,  
0.639344262295082,  
0.6557377049180327,  
0.6721311475409836,  
0.6885245901639344,  
0.7213114754098361,  
0.7049180327868853,  
0.6885245901639344,  
0.6885245901639344,  
0.7049180327868853,  
0.7540983606557377,  
0.7377049180327869,  
0.7377049180327869,  
0.7377049180327869,  
0.6885245901639344,  
0.7213114754098361,  
0.6885245901639344,  
0.6885245901639344,  
0.7049180327868853,  
0.6557377049180327]
```

```
[33]: plt.plot(neighbors,train_scores, label="Train score")  
plt.plot(neighbors,test_scores,label="Test score")  
plt.xticks(np.arange(1,21,1))  
plt.xlabel("Number of neighbors")  
plt.ylabel("Model score")  
plt.legend();  
  
print(f"Maximum KNN score on the test data: {max(test_scores)*100:.2f}%")
```

Maximum KNN score on the test data: 75.41%



[34]: #KNN is not the right algorithm for this problem because the other algorithms → perform far better

2.3 Hyperparameter tuning with RandomizedSearchCV

We are going to tune: * LogisticRegression() * RandomForestClassifier()

... using RSCV

```
[35]: # Create a hyperparameter grid for logistic regression model
log_reg_grid={
    "C":np.logspace(-4,4,20),
    "solver":["liblinear"]
}

# Create a HP grid fpr RandomForestClass.
rf_grid ={
    "n_estimators":np.arange(10,1000,50),
    "max_depth":[None,3,5,10],
    "min_samples_split": np.arange(2,20,2),
    "min_samples_leaf":np.arange(1,20,2)
}
```

Now we have got hyperparameters grids for each of our models, let's tune them using RandomizedSearchCV

```
[36]: # tune LR Model
```

```
np.random.seed(42)

# Setup random hyperparameter search for LogisticRegression
rs_log_reg = RandomizedSearchCV(LogisticRegression(),
                                 param_distributions=log_reg_grid,
                                 cv=5,
                                 n_iter = 20,
                                 verbose = True)

# Fit random hyperparameter search model for LogisticRegression
rs_log_reg.fit(X_train,y_train)
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits

```
[36]: RandomizedSearchCV(cv=5, estimator=LogisticRegression(), n_iter=20,
                        param_distributions={'C': array([1.0000000e-04,
                        2.63665090e-04, 6.95192796e-04, 1.83298071e-03,
                        4.83293024e-03, 1.27427499e-02, 3.35981829e-02, 8.85866790e-02,
                        2.33572147e-01, 6.15848211e-01, 1.62377674e+00, 4.28133240e+00,
                        1.12883789e+01, 2.97635144e+01, 7.84759970e+01, 2.06913808e+02,
                        5.45559478e+02, 1.43844989e+03, 3.79269019e+03, 1.0000000e+04]),
                        'solver': ['liblinear']},
                        verbose=True)
```

```
[37]: rs_log_reg.best_params_
```

```
[37]: {'solver': 'liblinear', 'C': np.float64(0.23357214690901212)}
```

```
[38]: rs_log_reg.score(X_test,y_test)
```

```
[38]: 0.8852459016393442
```

Now we have tuned LogisticRegression(), let's do the same for RandomForest()

```
[39]: np.random.seed(42)
```

```
# Setup random hyperparameter search for RandomForestClassifier
rs_rf = RandomizedSearchCV(RandomForestClassifier(),
                            param_distributions=rf_grid,
                            cv=5,
                            n_iter = 20,
                            verbose = True)

# Fit random hyperparameter search model for LogisticRegression
```

```
rs_rf.fit(X_train,y_train)
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits

```
[39]: RandomizedSearchCV(cv=5, estimator=RandomForestClassifier(), n_iter=20,
                         param_distributions={'max_depth': [None, 3, 5, 10],
                                              'min_samples_leaf': array([ 1,  3,  5,
                                              7,  9, 11, 13, 15, 17, 19]),
                                              'min_samples_split': array([ 2,  4,  6,
                                              8, 10, 12, 14, 16, 18]),
                                              'n_estimators': array([ 10,  60, 110,
                                              160, 210, 260, 310, 360, 410, 460, 510, 560, 610,
                                              660, 710, 760, 810, 860, 910, 960])},
                         verbose=True)
```

```
[40]: rs_rf.best_params_
```

```
[40]: {'n_estimators': np.int64(210),
       'min_samples_split': np.int64(4),
       'min_samples_leaf': np.int64(19),
       'max_depth': 3}
```

```
[41]: rs_rf.score(X_train,y_train)
```

```
[41]: 0.8553719008264463
```

2.4 Hyperparameter Tuning with GridSearchCV

Since our LogisticRegression model provides the best scores so far, we will try and improve the HP by GridSearchCV

```
[42]: # Diffrent hyperparameters for our LogisticRegression Model
log_reg_grid={
    "C":np.logspace(-4,4,30),
    "solver":["liblinear"]
}

# Setup grid hyperparameter search for LogisticRegression
gs_log_reg = GridSearchCV(LogisticRegression(),
                           param_grid=log_reg_grid,
                           cv = 5,
                           verbose = True)

# Fit the model
gs_log_reg.fit(X_train,y_train)
```

Fitting 5 folds for each of 30 candidates, totalling 150 fits

```
[42]: GridSearchCV(cv=5, estimator=LogisticRegression(),
                  param_grid={'C': array([1.0000000e-04, 1.88739182e-04,
```

```

3.56224789e-04, 6.72335754e-04,
    1.26896100e-03, 2.39502662e-03, 4.52035366e-03, 8.53167852e-03,
    1.61026203e-02, 3.03919538e-02, 5.73615251e-02, 1.08263673e-01,
    2.04335972e-01, 3.85662042e-01, 7.27895384e-01, 1.37382380e+00,
    2.59294380e+00, 4.89390092e+00, 9.23670857e+00, 1.74332882e+01,
    3.29034456e+01, 6.21016942e+01, 1.17210230e+02, 2.21221629e+02,
    4.17531894e+02, 7.88046282e+02, 1.48735211e+03, 2.80721620e+03,
    5.29831691e+03, 1.00000000e+04]),
                    'solver': ['liblinear']},
            verbose=True)

```

```
[43]: gs_log_reg.score(X_test,y_test)
```

[43]: 0.8852459016393442

2.5 Evaluating our tuned machine learning classifier, beyond accuracy

- ROC Curve and AUC Score
 - Confusun Matrix
 - Classification Report
 - Precision
 - Recall
 - F1-Score

... and it would be great if cross-validation was used where possible

To make comparisons and evaluate our trained model, first we need to make predictions.

```
[45]: # Make Predictions with tuned model  
y_preds = gs_log_reg.predict(X_test)
```

[46] : y_preds

```
[46]: array([0, 1, 1, 0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0,  
          0, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
          1, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0])
```

[47]: v_test

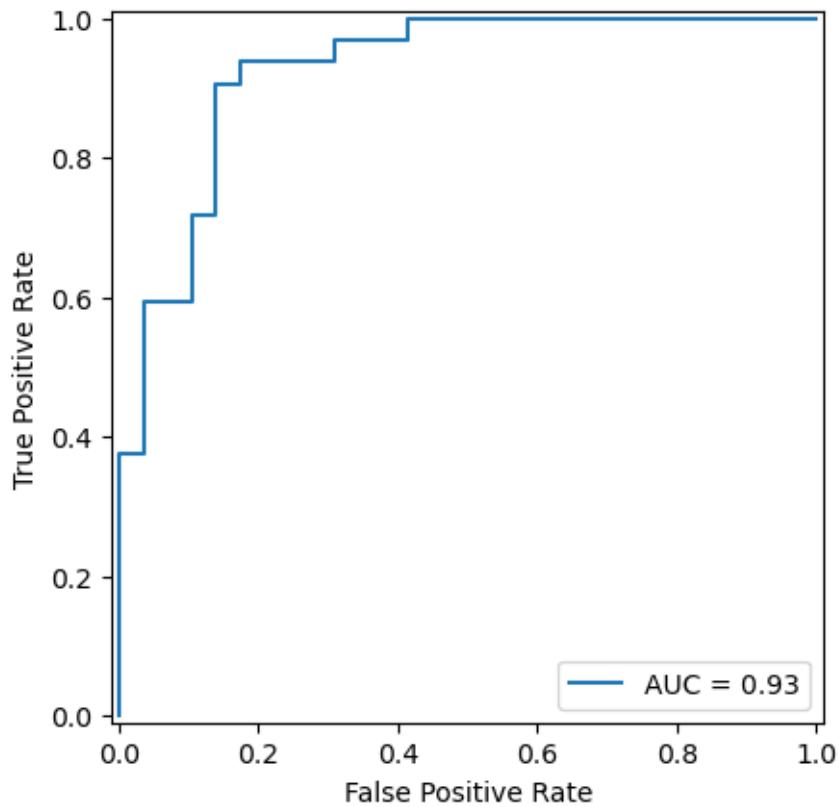
[47]:

179	0
228	0
111	1
246	0
60	1
..	
249	0
104	1
300	0
193	0

```
184      0  
Name: target, Length: 61, dtype: int64
```

```
[58]: from sklearn import metrics  
  
y_score = gs_log_reg.predict_proba(X_test)[:, 1]    # Wahrscheinlichkeit für  
        ↳ Klasse 1  
  
fpr, tpr, thresholds = metrics.roc_curve(y_test, y_score)  
roc_auc = metrics.auc(fpr, tpr)  
  
disp = metrics.RocCurveDisplay(fpr=fpr, tpr=tpr, roc_auc=roc_auc)  
disp.plot()
```

```
[58]: <sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x223bc39f790>
```



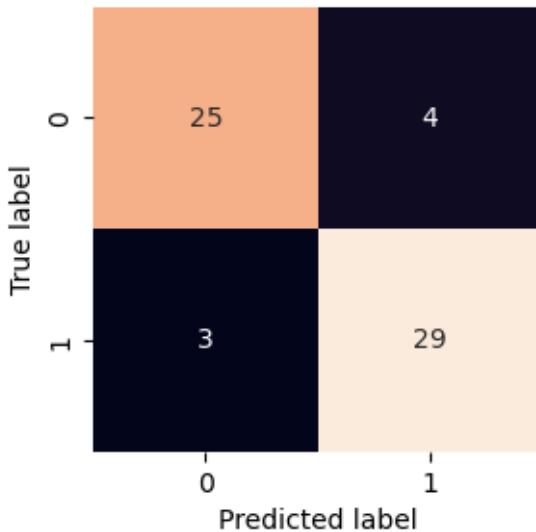
```
[59]: y_score
```

```
[59]: array([0.13274739, 0.75591518, 0.81452426, 0.05469225, 0.88453364,  
        0.87070287, 0.60512182, 0.00435981, 0.01376378, 0.56138489,  
        0.7172079 , 0.11904138, 0.88730101, 0.06005038, 0.96750057,
```

```
0.93181128, 0.96404698, 0.09452756, 0.01769764, 0.0264636 ,  
0.71543156, 0.02727603, 0.14274667, 0.71660794, 0.88198278,  
0.69480581, 0.8423425 , 0.69335569, 0.01830941, 0.87782365,  
0.07150765, 0.06684154, 0.01510285, 0.14314483, 0.60229465,  
0.12640328, 0.6633502 , 0.85079097, 0.81898326, 0.84121548,  
0.54515856, 0.79250831, 0.77817602, 0.70538842, 0.83243213,  
0.02113004, 0.73216369, 0.93234387, 0.10276671, 0.06440756,  
0.13470669, 0.03554564, 0.80441973, 0.95312794, 0.31714658,  
0.00309601, 0.08734727, 0.93823575, 0.02813797, 0.01309874,  
0.06291747])
```

```
[63]: gs_log_reg.best_params_
```

```
def plot_conf_mat(y_test, y_preds):  
    """  
    Plots a confusion matrix using Seaborn's heatmap().  
    """  
    fig, ax = plt.subplots(figsize=(3, 3))  
    ax = sns.heatmap(confusion_matrix(y_test, y_preds),  
                      annot=True, # Annotate the boxes  
                      cbar=False)  
    plt.xlabel("Predicted label") # predictions go on the x-axis  
    plt.ylabel("True label") # true labels go on the y-axis  
  
plot_conf_mat(y_test, y_preds)
```



```
[61]: confusion_matrix(y_test, y_preds)
```

```
[61]: array([[25,  4],  
           [ 3, 29]])
```

Now we've got a ROC curve, an AUC metric and a confusion matrix, let's get a classification report as well as cross-validated precision, recall and f1-score

```
[64]: print(classification_report(y_test,y_preds))
```

	precision	recall	f1-score	support
0	0.89	0.86	0.88	29
1	0.88	0.91	0.89	32
accuracy			0.89	61
macro avg	0.89	0.88	0.88	61
weighted avg	0.89	0.89	0.89	61

2.5.1 Calculating evaluation metrics using cross-validation

We're going to calculate precision, recall and f1-score of our model using cross-validation and to do so we'll be using `cross_val_score()`

```
[65]: # Check best Hyperparameters  
gs_log_reg.best_params_
```

```
[65]: {'C': np.float64(0.20433597178569418), 'solver': 'liblinear'}
```

```
[66]: # Create a new classifier with best parameters  
clf = LogisticRegression(C= 0.20433597178569418, solver = "liblinear")
```

```
[73]: # Cross-validated accuracy  
cv_acc = cross_val_score(clf,X,y,cv = 5, scoring = "accuracy")  
  
cv_acc.mean()
```

```
[73]: np.float64(0.8446994535519124)
```

```
[74]: #Cross-validated precision  
cv_prec = cross_val_score(clf,X,y,cv = 5, scoring = "precision")  
cv_prec.mean()
```

```
[74]: np.float64(0.8207936507936507)
```

```
[75]: #Cross-validated recall  
cv_rec = cross_val_score(clf,X,y,cv = 5, scoring = "recall")  
cv_rec.mean()
```

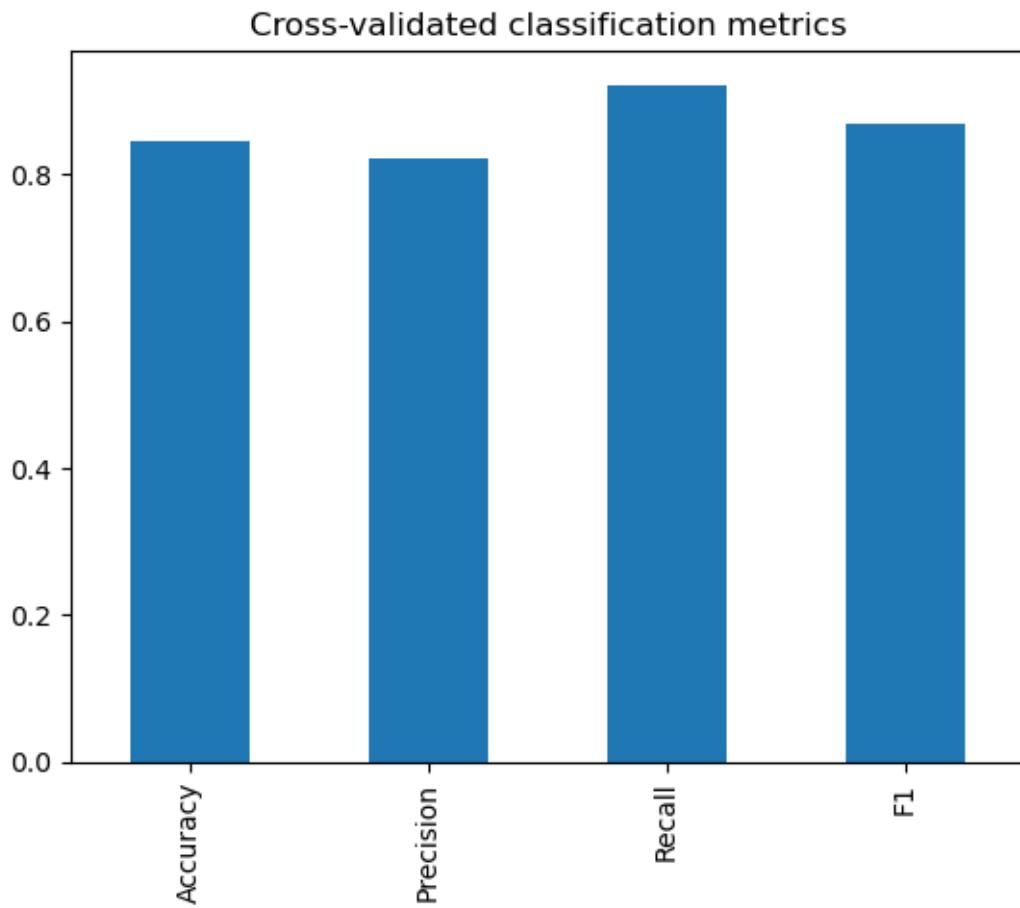
```
[75]: np.float64(0.9212121212121213)
```

```
[78]: #Cross-validates f1-score
cv_f1 = cross_val_score(clf,X,y,cv = 5, scoring = "f1")
cv_f1.mean()

[78]: np.float64(0.8673007976269721)

[85]: # Visualize cross-validated metrics
cv_metrics = pd.DataFrame({"Accuracy": cv_acc.mean(),
                            "Precision":cv_prec.mean(),
                            "Recall":cv_rec.mean(),
                            "F1":cv_f1.mean()},
                           index = [0])
cv_metrics.T.plot.bar(title="Cross-validated classification metrics", legend=False)

[85]: <Axes: title={'center': 'Cross-validated classification metrics'}>
```



2.5.2 Feature Importance

Feature importance is another way of asking which features contributed most to the outcomes of the model and how did they contribute

Finding feature importance is different for each machine learning model.

Let's find the feature importance for our LR model...

```
[86]: df.head()
```

```
[86]:    age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  slope \
0    63    1    3      145    233    1      0     150      0      2.3      0
1    37    1    2      130    250    0      1     187      0      3.5      0
2    41    0    1      130    204    0      0     172      0      1.4      2
3    56    1    1      120    236    0      1     178      0      0.8      2
4    57    0    0      120    354    0      1     163      1      0.6      2

      ca  thal  target
0    0    1      1
1    0    2      1
2    0    2      1
3    0    2      1
4    0    2      1
```

```
[88]: # Fit an instance of Logistic Regression
# Create a new classifier with best parameters
clf = LogisticRegression(C= 0.20433597178569418, solver = "liblinear")

clf.fit(X_train,y_train);
```

```
[89]: #Check coef_
clf.coef_
```

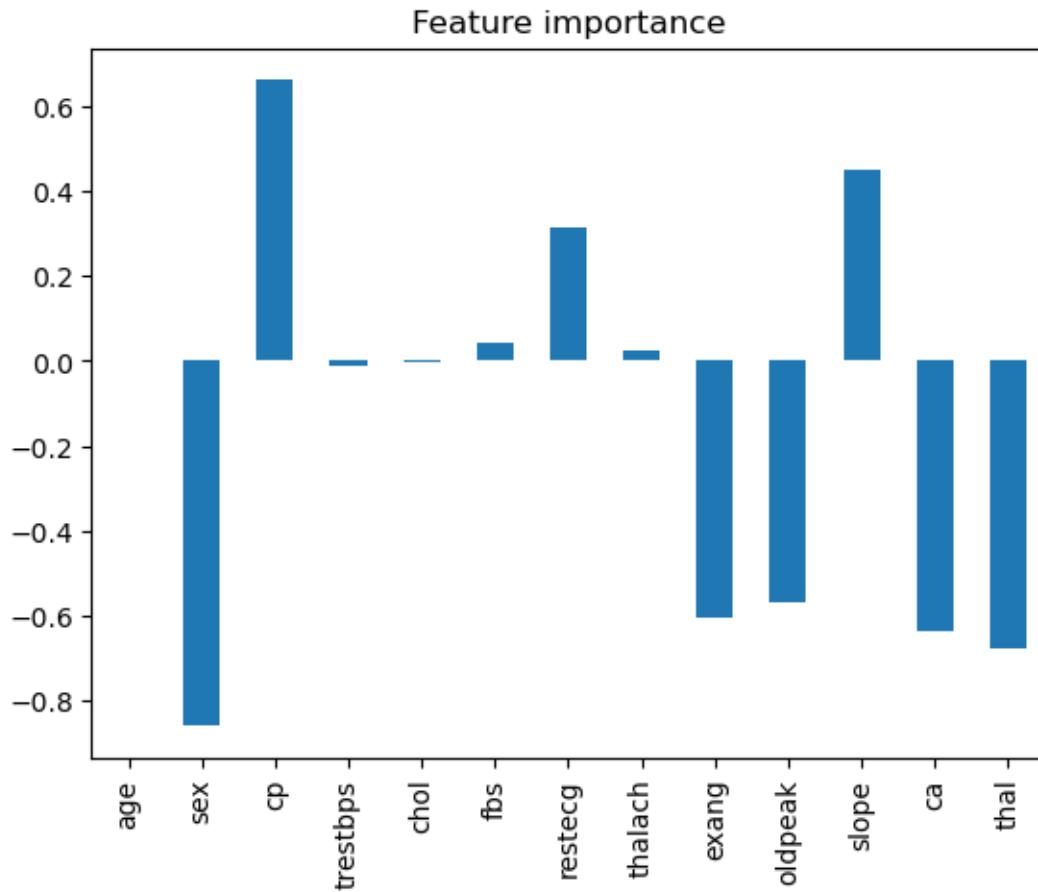
```
[89]: array([[ 0.00316727, -0.86044582,   0.66067073, -0.01156993, -0.00166374,
       0.04386131,   0.31275787,   0.02459361, -0.60413038, -0.56862852,
       0.45051617, -0.63609863, -0.67663375]])
```

```
[90]: feature_dict = dict(zip(df.columns,list(clf.coef_[0])))
feature_dict
```

```
[90]: {'age': np.float64(0.0031672721856887734),
       'sex': np.float64(-0.860445816920919),
       'cp': np.float64(0.6606707303492849),
       'trestbps': np.float64(-0.011569930902919925),
       'chol': np.float64(-0.001663741604035976),
       'fbs': np.float64(0.04386130751482091),
       'restecg': np.float64(0.3127578715206996),
       'thalach': np.float64(0.02459360818122666),
```

```
'exang': np.float64(-0.6041303799858143),
'oldpeak': np.float64(-0.5686285194546157),
'slope': np.float64(0.4505161679452401),
'ca': np.float64(-0.6360986316921434),
'thal': np.float64(-0.6766337521354281)}
```

```
[91]: # Visualize Feature Importance
feature_df = pd.DataFrame(feature_dict, index =[0])
feature_df.T.plot.bar(title="Feature importance", legend = False);
```



```
[92]: pd.crosstab(df["sex"], df["target"])
```

```
[92]: target      0      1
      sex
      0        24    72
      1       114   93
```

```
[94]: pd.crosstab(df["slope"], df["target"]) # slope: the slope of the peak exercise
      ↳ ST segment
```

```
[94]: target    0    1  
slope  
0      12    9  
1      91   49  
2      35  107
```

2.6 6. Experiments

If you haven't hit your evaluation metric yet... ask yourself...

- Could you collect more data?
- Could you try a better model? Like CatBoost or XGBoost?
- Could you improve the current models? (beyond what we've done so far)
- If your model is good enough (you have hit your evaluation metric)
- How would you export it and share it

[]: